

ATM Case Study, Part I: Object-Oriented Design with the UML

30

*Action speaks louder than words
but not nearly as often.*

—Mark Twain

*Always design a thing by
considering it in its next larger
context.*

—Eliel Saarinen

Oh, life is a glorious cycle of song.

—Dorothy Parker

*The Wright brothers' design ...
allowed them to survive long
enough to learn how to fly.*

—Michael Potts

Objectives

In this chapter you'll learn:

- A simple object-oriented design methodology.
- What a requirements document is.
- To identify classes and class attributes from a requirements document.
- To identify objects' states, activities and operations from a requirements document.
- To determine the collaborations among objects in a system.
- To work with various UML diagrams to graphically model an object-oriented system.

30.1 Introduction	30.5 Identifying Objects' States and Activities
30.2 Examining the ATM Requirements Document	30.6 Identifying Class Operations
30.3 Identifying the Classes in the ATM Requirements Document	30.7 Identifying Collaboration Among Objects
30.4 Identifying Class Attributes	30.8 Wrap-Up

Answers to Self-Review Exercises

30.1 Introduction

Now we begin the optional portion of our object-oriented design and implementation case study. In this chapter and Chapter 31, you'll design and implement an object-oriented automated teller machine (ATM) software system. The case study provides you with a concise, carefully paced, complete design and implementation experience. In Sections 30.2–30.7 and 31.2–31.3, you'll perform the steps of an object-oriented design (OOD) process using the UML while relating these steps to the concepts discussed in Chapters 3–12. In this chapter, you'll work with six popular types of UML diagrams to graphically represent the design. In Chapter 31, you'll tune the design with inheritance, then fully implement the ATM in a C# console application (Section 31.4).

This is not an exercise; rather, it's an end-to-end learning experience that concludes with a detailed walkthrough of the complete C# code that implements our design. It will acquaint you with the kinds of substantial problems encountered in industry. These chapters can be studied as a continuous unit after you've completed the introduction to object-oriented programming in Chapters 3 and 10–12.

30.2 Examining the ATM Requirements Document

We begin our design process by presenting a **requirements document** that specifies the overall purpose of the ATM system and *what* it must do. Throughout the case study, we refer to the requirements document to determine precisely what functionality the system must include.

Requirements Document

A small local bank intends to install a new automated teller machine (ATM) to allow users (i.e., bank customers) to perform basic financial transactions (Fig. 30.1). For simplicity, each user can have only one account at the bank. ATM users should be able to view their account balance, withdraw cash (i.e., take money out of an account) and deposit funds (i.e., place money into an account).

The user interface of the automated teller machine contains the following hardware components:

- a screen that displays messages to the user
- a keypad that receives numeric input from the user
- a cash dispenser that dispenses cash to the user
- a deposit slot that receives deposit envelopes from the user

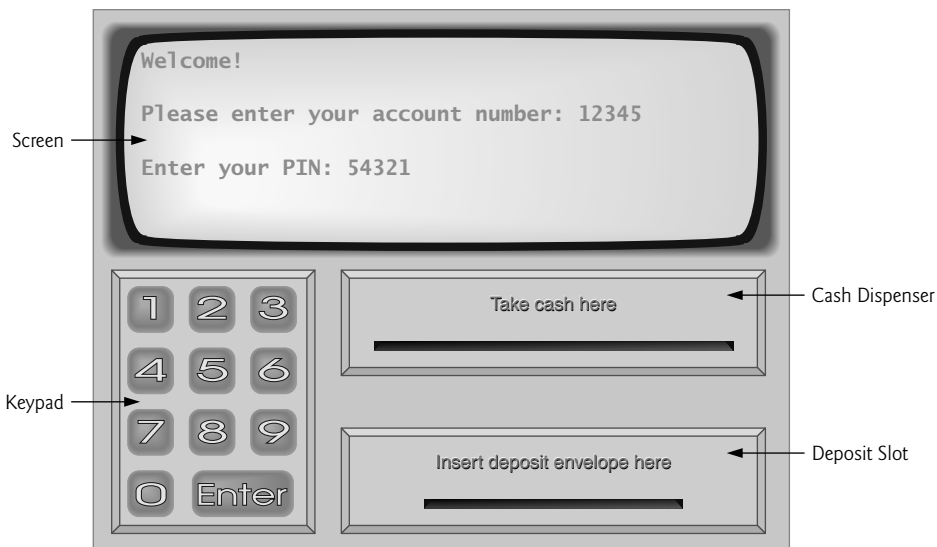


Fig. 30.1 | Automated teller machine user interface.

The cash dispenser begins each day loaded with 500 \$20 bills. [Note: Owing to the limited scope of this case study, certain elements of the ATM described here simplify various aspects of a real ATM. For example, a real ATM typically contains a device that reads a user's account number from an ATM card, whereas this ATM asks the user to type an account number on the keypad (which you'll simulate with your personal computer's keypad). Also, a real ATM usually prints a paper receipt at the end of a session, but all output from this ATM appears on the screen.]

The bank wants you to develop software to perform the financial transactions initiated by bank customers through the ATM. The bank will integrate the software with the ATM's hardware at a later time. The software should simulate the functionality of the hardware devices (e.g., cash dispenser, deposit slot) in software components, but it need not concern itself with how these devices perform their duties. The ATM hardware has not been developed yet, so instead of writing your software to run on the ATM, you should develop a first version of the software to run on a personal computer. This version should use the computer's monitor to simulate the ATM's screen and the computer's keyboard to simulate the ATM's keypad.

An ATM session consists of authenticating a user (i.e., proving the user's identity) based on an account number and personal identification number (PIN), followed by creating and executing financial transactions. To authenticate a user and perform transactions, the ATM must interact with the bank's account information database. [Note: A database is an organized collection of data stored on a computer.] For each bank account, the database stores an account number, a PIN and a balance indicating the amount of money in the account. [Note: The bank plans to build only one ATM, so we do not need to worry about multiple ATMs accessing the database at the same time. Furthermore, we assume that the bank does not make any changes to the information in the database while

a user is accessing the ATM. Also, any business system like an ATM faces reasonably complicated security issues that go well beyond the scope of a first- or second-semester programming course. We make the simplifying assumption, however, that the bank trusts the ATM to access and manipulate the information in the database without significant security measures.]

Upon approaching the ATM, the user should experience the following sequence of events (see Fig. 30.1):

1. The screen displays a welcome message and prompts the user to enter an account number.
2. The user enters a five-digit account number, using the keypad.
3. For authentication purposes, the screen prompts the user to enter the PIN (personal identification number) associated with the specified account number.
4. The user enters a five-digit PIN, using the keypad.
5. If the user enters a valid account number and the correct PIN for that account, the screen displays the main menu (Fig. 30.2). If the user enters an invalid account number or an incorrect PIN, the screen displays an appropriate message, then the ATM returns to *Step 1* to restart the authentication process.

After the ATM authenticates the user, the main menu (Fig. 30.2) displays a numbered option for each of the three types of transactions: balance inquiry (option 1), withdrawal (option 2) and deposit (option 3). The main menu also displays an option that allows the user to exit the system (option 4). The user then chooses either to perform a transaction (by entering 1, 2 or 3) or to exit the system (by entering 4). If the user enters an invalid option, the screen displays an error message, then redisplay the main menu.

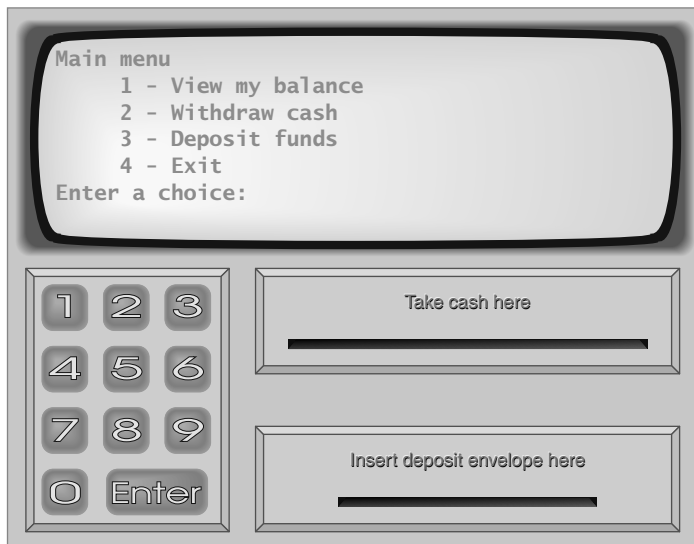


Fig. 30.2 | ATM main menu.

If the user enters 1 to make a balance inquiry, the screen displays the user's account balance. To do so, the ATM must retrieve the balance from the bank's database.

The following actions occur when the user enters 2 to make a withdrawal:

1. The screen displays a menu (shown in Fig. 30.3) containing standard withdrawal amounts: \$20 (option 1), \$40 (option 2), \$60 (option 3), \$100 (option 4) and \$200 (option 5). The menu also contains option 6, which allows the user to cancel the transaction.

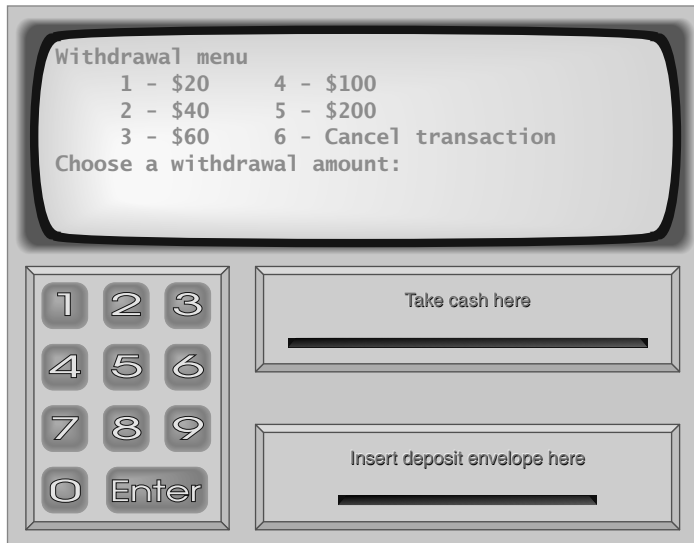


Fig. 30.3 | ATM withdrawal menu.

2. The user enters a menu selection (1–6) using the keypad.
3. If the withdrawal amount chosen is greater than the user's account balance, the screen displays a message stating this and telling the user to choose a smaller amount. The ATM then returns to *Step 1*. If the withdrawal amount chosen is less than or equal to the user's account balance (i.e., an acceptable withdrawal amount), the ATM proceeds to *Step 4*. If the user chooses to cancel the transaction (option 6), the ATM displays the main menu (Fig. 30.2) and waits for user input.
4. If the cash dispenser contains enough cash to satisfy the request, the ATM proceeds to *Step 5*. Otherwise, the screen displays a message indicating the problem and telling the user to choose a smaller withdrawal amount. The ATM then returns to *Step 1*.
5. The ATM debits (i.e., subtracts) the withdrawal amount from the user's account balance in the bank's database.
6. The cash dispenser dispenses the desired amount of money to the user.
7. The screen displays a message reminding the user to take the money.

The following actions occur when the user enters 3 (from the main menu) to make a deposit:

1. The screen prompts the user to enter a deposit amount or to type 0 (zero) to cancel the transaction.
2. The user enters a deposit amount or 0, using the keypad. [*Note:* The keypad does not contain a decimal point or a dollar sign, so the user cannot type a real dollar amount (e.g., \$147.25). Instead, the user must enter a deposit amount as a number of cents (e.g., 14725). The ATM then divides this number by 100 to obtain a number representing a dollar amount (e.g., $14725 \div 100 = 147.25$).]
3. If the user specifies a deposit amount, the ATM proceeds to *Step 4*. If the user chooses to cancel the transaction (by entering 0), the ATM displays the main menu (Fig. 30.2) and waits for user input.
4. The screen displays a message telling the user to insert a deposit envelope into the deposit slot.
5. If the deposit slot receives a deposit envelope within two minutes, the ATM credits (i.e., adds) the deposit amount to the user's account balance in the bank's database. [*Note:* This money is not immediately available for withdrawal. The bank first must verify the amount of cash in the deposit envelope, and any checks in the envelope must clear (i.e., money must be transferred from the check writer's account to the check recipient's account). When either of these events occurs, the bank appropriately updates the user's balance stored in its database. This occurs independently of the ATM system.] If the deposit slot does not receive a deposit envelope within two minutes, the screen displays a message that the system has canceled the transaction due to inactivity. The ATM then displays the main menu and waits for user input.

After the system successfully executes a transaction, the system should redisplay the main menu (Fig. 30.2) so that the user can perform additional transactions. If the user chooses to exit the system (by entering option 4), the screen should display a thank-you message, then display the welcome message for the next user.

Analyzing the ATM System

The preceding statement presented a simplified requirements document. Typically, such a document is the result of a detailed process of **requirements gathering** that might include interviews with potential users of the system and specialists in fields related to the system. For example, a systems analyst who is hired to prepare a requirements document for banking software (e.g., the ATM system described here) might interview financial experts and people who have used ATMs to gain a better understanding of *what* the software must do. The analyst would use the information gained to compile a list of **system requirements** to guide systems designers.

The process of requirements gathering is a key task of the first stage of the software life cycle. The **software life cycle** specifies the stages through which software evolves from the time it's conceived to the time it's retired from use. These stages typically include analysis, design, implementation, testing and debugging, deployment, maintenance and retirement. Several software life-cycle models exist, each with its own preferences and

specifications for when and how often software engineers should perform the various stages. **Waterfall models** perform each stage once in succession, whereas **iterative models** may repeat one or more stages several times throughout a product's life cycle.

The analysis stage of the software life cycle focuses on precisely defining the problem to be solved. When designing any system, one must certainly *solve the problem right*, but of equal importance, one must *solve the right problem*. Systems analysts collect the requirements that indicate the specific problem to solve. Our requirements document describes our simple ATM system in sufficient detail that you do not need to go through an extensive analysis stage—it has been done for you.

To capture what a proposed system should do, developers often employ a technique known as **use case modeling**. This process identifies the **use cases** of the system, each of which represents a different capability that the system provides to its clients. For example, ATMs typically have several use cases, such as “View Account Balance,” “Withdraw Cash,” “Deposit Funds,” “Transfer Funds Between Accounts” and “Buy Postage Stamps.” The simplified ATM system we build in this case study requires only the first three use cases (Fig. 30.4).




Fig. 30.4 | Use case diagram for the ATM system from the user's perspective.

Each use case describes a typical scenario in which the user uses the system. You have already read descriptions of the ATM system's use cases in the requirements document; the lists of steps required to perform each type of transaction (i.e., balance inquiry, withdrawal and deposit) actually described the three use cases of our ATM—“View Account Balance,” “Withdraw Cash” and “Deposit Funds.”

Use Case Diagrams

We now introduce the first of several UML diagrams in our ATM case study. We create a **use case diagram** to model the interactions between a system's clients (in this case study, bank customers) and the system. The goal is to show the kinds of interactions users have with a system without providing the details—these are shown in other UML diagrams (which we present throughout the case study). Use case diagrams are often accompanied by informal text that describes the use cases in more detail—like the text that appears in the requirements document. Use case diagrams are produced during the analysis stage of the software life cycle. In larger systems, use case diagrams are simple but indispensable tools that help system designers focus on satisfying the users' needs.

Figure 30.4 shows the use case diagram for our ATM system. The stick figure represents an **actor**, which defines the roles that an external entity—such as a person or another system—plays when interacting with the system. For our automated teller machine, the actor is a User who can view an account balance, withdraw cash and deposit funds using the ATM. The User is not an actual person, but instead comprises the roles that a real person—when playing the part of a User—can play while interacting with the ATM. A use case diagram can include multiple actors. For example, the use case diagram for a real bank's ATM system might also include an actor named Administrator who refills the cash dispenser each day.

We identify the actor in our system by examining the requirements document, which states, “ATM users should be able to view their account balance, withdraw cash and deposit funds.” The actor in each of the three use cases is simply the User who interacts with the ATM. An external entity—a real person—plays the part of the User to perform financial transactions. Figure 30.4 shows one actor, whose name, User, appears below the actor in the diagram. The UML models each use case as an oval connected to an actor with a solid line.

Software engineers (more precisely, systems designers) must analyze the requirements document, or a set of use cases, and design the system before programmers implement it in a particular programming language. During the analysis stage, systems designers focus on understanding the requirements document to produce a high-level specification that describes *what* the system is supposed to do. The output of the design stage—a **design specification**—should specify *how* the system should be constructed to satisfy these requirements. In the next several Software Engineering Case Study sections, we perform the steps of a simple OOD process on the ATM system to produce a design specification containing a collection of UML diagrams and supporting text. Recall that the UML is designed for use with any OOD process. Many such processes exist, the best known being the Rational Unified Process™ (RUP) developed by Rational Software Corporation (now a division of IBM). RUP is a rich process for designing “industrial-strength” applications. For this case study, we present a simplified design process.

Designing the ATM System

We now begin the design stage of our ATM system. A **system** is a set of components that interact to solve a problem. For example, to perform the ATM system's designated tasks, our ATM system has a user interface (Fig. 30.1), contains software that executes financial transactions and interacts with a database of bank-account information. **System structure** describes the system's objects and their interrelationships. **System behavior** describes how the system changes as its objects interact with one another. Every system has both structure and behavior—designers must specify both. There are several distinct types of system structures and behaviors. For example, the interactions among objects in the system differ from those between the user and the system, yet both constitute a portion of the system behavior.

The UML 2 specifies 13 diagram types for documenting system models. Each diagram type models a distinct characteristic of a system's structure or behavior—six relate to system structure and seven to system behavior. We list here only the six types of diagrams used in our case study—of which one (the class diagram) models system structure and five model system behavior. We overview the remaining seven UML diagram types in Appendix E, UML 2: Additional Diagram Types.

1. **Use case diagrams**, such as the one in Fig. 30.4, model the interactions between a system and its external entities (actors) in terms of use cases (system capabilities, such as “View Account Balance,” “Withdraw Cash” and “Deposit Funds”).
2. **Class diagrams**, which you’ll study in Section 30.3, model the classes, or “building blocks,” used in a system. Each noun, or “thing,” described in the requirements document is a candidate to be a class in the system (e.g., “account,” “keypad”). Class diagrams help us specify the structural relationships between parts of the system. For example, the ATM system class diagram will, among other things, specify that the ATM is physically composed of a screen, a keypad, a cash dispenser and a deposit slot.
3. **State machine diagrams**, which you’ll study in Section 30.5, model the ways in which an object changes state. An object’s **state** is indicated by the values of all its attributes at a given time. When an object changes state, it may subsequently behave differently in the system. For example, after validating a user’s PIN, the ATM transitions from the “user not authenticated” state to the “user authenticated” state, at which point the ATM allows the user to perform financial transactions (e.g., view account balance, withdraw cash, deposit funds).
4. **Activity diagrams**, which you’ll also study in Section 30.5, model an object’s **activity**—the object’s workflow (sequence of events) during program execution. An activity diagram models the actions the object performs and specifies the order in which it performs them. For example, an activity diagram shows that the ATM must obtain the balance of the user’s account (from the bank’s account-information database) before the screen can display the balance to the user.
5. **Communication diagrams** (called collaboration diagrams in earlier versions of the UML) model the interactions among objects in a system, with an emphasis on *what* interactions occur. You’ll learn in Section 30.7 that these diagrams show which objects must interact to perform an ATM transaction. For example, the ATM must communicate with the bank’s account-information database to retrieve an account balance.
6. **Sequence diagrams** also model the interactions among the objects in a system, but unlike communication diagrams, they emphasize *when* interactions occur. You’ll learn in Section 30.7 that these diagrams help show the order in which interactions occur in executing a financial transaction. For example, the screen prompts the user to enter a withdrawal amount before cash is dispensed.

In Section 30.3, we continue designing our ATM system by identifying the classes from the requirements document. We accomplish this by extracting key nouns and noun phrases from the requirements document. Using these classes, we develop our first draft of the class diagram that models the structure of our ATM system.

Web Resources

We’ve created an extensive UML Resource Center that contains many links to additional information, including introductions, tutorials, blogs, books, certification, conferences, developer tools, documentation, e-books, FAQs, forums, groups, UML in Java, podcasts,

security, tools, downloads, training courses, videos and more. We encourage you to browse our UML Resource Center at www.deitel.com/UML/ to learn more.

Self-Review Exercises

30.1 Suppose we enabled a user of our ATM system to transfer money between two bank accounts. Modify the use case diagram of Fig. 30.4 to reflect this change.

30.2 _____ model the interactions among objects in a system with an emphasis on *when* these interactions occur.

- a) Class diagrams
- b) Sequence diagrams
- c) Communication diagrams
- d) Activity diagrams

30.3 Which of the following choices lists stages of a typical software life cycle in sequential order?

- a) design, analysis, implementation, testing
- b) design, analysis, testing, implementation
- c) analysis, design, testing, implementation
- d) analysis, design, implementation, testing

30.3 Identifying the Classes in the ATM Requirements Document

Now we begin designing the ATM system. In this section, we identify the classes that are needed to build the ATM system by analyzing the nouns and noun phrases that appear in the requirements document. We introduce UML class diagrams to model the relationships between these classes. This is an important first step in defining the structure of our system.

Identifying the Classes in a System

We begin our object-oriented design (OOD) process by identifying the classes required to build the ATM system. We'll eventually describe these classes using UML class diagrams and implement these classes in C#. First, we review the requirements document of Section 30.2 and find key nouns and noun phrases to help us identify classes that comprise the ATM system. We may decide that some of these nouns and noun phrases are attributes of other classes in the system. We may also conclude that some of the nouns and noun phrases do not correspond to parts of the system and thus should not be modeled at all. Additional classes may become apparent to us as we proceed through the design process. Figure 30.5 lists the nouns and noun phrases in the requirements document.

Nouns and noun phrases in the requirements document		
bank	money / funds	account number
ATM	screen	PIN
user	keypad	bank database
customer	cash dispenser	balance inquiry

Fig. 30.5 | Nouns and noun phrases in the requirements document.

Nouns and noun phrases in the requirements document		
transaction	\$20 bill / cash	withdrawal
account	deposit slot	deposit
balance	deposit envelope	

Fig. 30.5 | Nouns and noun phrases in the requirements document.

We create classes only for the nouns and noun phrases that have significance in the ATM system. We do not need to model “bank” as a class, because the bank is not a part of the ATM system—the bank simply wants us to build the ATM. “User” and “customer” also represent entities outside of the system—they’re important because they interact with our ATM system, but we do not need to model them as classes in the ATM system. Recall that we modeled an ATM user (i.e., a bank customer) as the actor in the use case diagram of Fig. 30.4.

We do not model “\$20 bill” or “deposit envelope” as classes. These are physical objects in the real world, but they’re not part of what is being automated. We can adequately represent the presence of bills in the system using an attribute of the class that models the cash dispenser. (We assign attributes to classes in Section 30.4.) For example, the cash dispenser maintains a count of the number of bills it contains. The requirements document does not say anything about what the system should do with deposit envelopes after it receives them. We can assume that simply acknowledging the receipt of an envelope—an **operation** performed by the class that models the deposit slot—is sufficient to represent the presence of an envelope in the system. (We assign operations to classes in Section 30.6.)

In our simplified ATM system, representing various amounts of “money,” including the “balance” of an account, as attributes of other classes seems most appropriate. Likewise, the nouns “account number” and “PIN” represent significant pieces of information in the ATM system. They’re important attributes of a bank account. They do not, however, exhibit behaviors. Thus, we can most appropriately model them as attributes of an account class.

Though the requirements document frequently describes a “transaction” in a general sense, we do not model the broad notion of a financial transaction at this time. Instead, we model the three types of transactions (i.e., “balance inquiry,” “withdrawal” and “deposit”) as individual classes. These classes possess specific attributes needed to execute the transactions they represent. For example, a withdrawal needs to know the amount of money the user wants to withdraw. A balance inquiry, however, does not require any additional data if the user is authenticated. Furthermore, the three transaction classes exhibit unique behaviors. A withdrawal involves dispensing cash to the user, whereas a deposit involves receiving a deposit envelope from the user. [Note: In Section 31.3, we “factor out” common features of all transactions into a general “transaction” class using the object-oriented concepts of abstract classes and inheritance.]

We determine the classes for our system based on the remaining nouns and noun phrases from Fig. 30.5. Each of these refers to one or more of the following:

- ATM

- screen
- keypad
- cash dispenser
- deposit slot
- account
- bank database
- balance inquiry
- withdrawal
- deposit

The elements of this list are likely to be classes we'll need to implement our system, although it's too early in our design process to claim that this list is complete.

We can now model our system's classes based on the list we've created. We capitalize class names in the design process—a UML convention—as we'll do when we write the C# code that implements our design. If the name of a class contains more than one word, we run the words together and capitalize each word (e.g., `MultipleWordName`). Using these conventions, we create classes `ATM`, `Screen`, `Keypad`, `CashDispenser`, `DepositSlot`, `Account`, `BankDatabase`, `BalanceInquiry`, `Withdrawal` and `Deposit`. We construct our system using all of these classes as building blocks. Before we begin building the system, however, we must gain a better understanding of how the classes relate to one another.

Modeling Classes

The UML enables us to model, via **class diagrams**, the classes in the ATM system and their interrelationships. Figure 30.6 represents class `ATM`. In the UML, each class is modeled as a rectangle with three compartments. The top compartment contains the name of the class, centered horizontally and appearing in boldface. The middle compartment contains the class's attributes. (We discuss attributes in Sections 30.4–30.5.) The bottom compartment contains the class's operations (discussed in Section 30.6). In Fig. 30.6, the middle and bottom compartments are empty, because we've not yet determined this class's attributes and operations.



Fig. 30.6 | Representing a class in the UML using a class diagram.

Class diagrams also show the relationships between the classes of the system. Figure 30.7 shows how our classes `ATM` and `Withdrawal` relate to one another. For the moment, we choose to model only this subset of the ATM classes for simplicity. We present a more complete class diagram later in this section. Notice that the rectangles representing classes in this diagram are not subdivided into compartments. The UML allows the suppression of class attributes and operations in this manner, when appropriate, to create more

readable diagrams. Such a diagram is said to be an **elided diagram**—one in which some information, such as the contents of the second and third compartments, is not modeled. We'll place information in these compartments in Sections 30.4–30.6.




Fig. 30.7 | Class diagram showing an association among classes.

In Fig. 30.7, the solid line that connects the two classes represents an **association**—a relationship between classes. The numbers near each end of the line are **multiplicity** values, which indicate how many objects of each class participate in the association. In this case, following the line from one end to the other reveals that, at any given moment, one ATM object participates in an association with either zero or one Withdrawal objects—zero if the current user is not performing a transaction or has requested a different type of transaction, and one if the user has requested a withdrawal. The UML can model many types of multiplicity. Figure 30.8 explains the multiplicity types.

Symbol	Meaning
0	None
1	One
<i>m</i>	An integer value
0..1	Zero or one
<i>m</i> , <i>n</i>	<i>m</i> or <i>n</i>
<i>m</i> .. <i>n</i>	At least <i>m</i> , but not more than <i>n</i>
*	Any nonnegative integer (zero or more)
0..*	Zero or more (identical to *)
1..*	One or more

Fig. 30.8 | Multiplicity types.

An association can be named. For example, the word *Executes* above the line connecting classes ATM and Withdrawal in Fig. 30.7 indicates the name of that association. This part of the diagram reads “one object of class ATM executes zero or one objects of class Withdrawal.” association names are directional, as indicated by the filled arrowhead—so it would be improper, for example, to read the preceding association from right to left as “zero or one objects of class Withdrawal execute one object of class ATM.”

The word *currentTransaction* at the Withdrawal end of the association line in Fig. 30.7 is a **role name**, which identifies the role the Withdrawal object plays in its relationship with the ATM. A role name adds meaning to an association between classes by identifying the role a class plays in the context of an association. A class can play several roles in the same system. For example, in a college personnel system, a person may play the role of “professor” when relating to students. The same person may take on the role of “col-

league” when participating in a relationship with another professor, and “coach” when coaching student athletes. In Fig. 30.7, the role name `currentTransaction` indicates that the `Withdrawal` object participating in the `Executes` association with an object of class `ATM` represents the transaction currently being processed by the ATM. In other contexts, a `Withdrawal` object may take on other roles (e.g., the previous transaction). Notice that we do not specify a role name for the `ATM` end of the `Executes` association. Role names are often omitted in class diagrams when the meaning of an association is clear without them.

In addition to indicating simple relationships, associations can specify more complex relationships, such as objects of one class being composed of objects of other classes. Consider a real-world automated teller machine. What “pieces” does a manufacturer put together to build a working ATM? Our requirements document tells us that the ATM is composed of a screen, a keypad, a cash dispenser and a deposit slot.

In Fig. 30.9, the **solid diamonds** attached to the association lines of class `ATM` indicate that class `ATM` has a **composition** relationship with classes `Screen`, `Keypad`, `CashDispenser` and `DepositSlot`. Composition implies a whole/part relationship. The class that has the composition symbol (the solid diamond) on its end of the association line is the whole (in this case, `ATM`), and the classes on the other end of the association lines are the parts—in this case, classes `Screen`, `Keypad`, `CashDispenser` and `DepositSlot`. The compositions in Fig. 30.9 indicate that an object of class `ATM` is formed from one object of class `Screen`, one object of class `CashDispenser`, one object of class `Keypad` and one object of class `DepositSlot`—the ATM “has a” screen, a keypad, a cash dispenser and a deposit slot. The **has-a relationship** defines composition. (We’ll see in the Software Engineering Case Study section in Section 31.3 that the *is-a* relationship defines inheritance.)

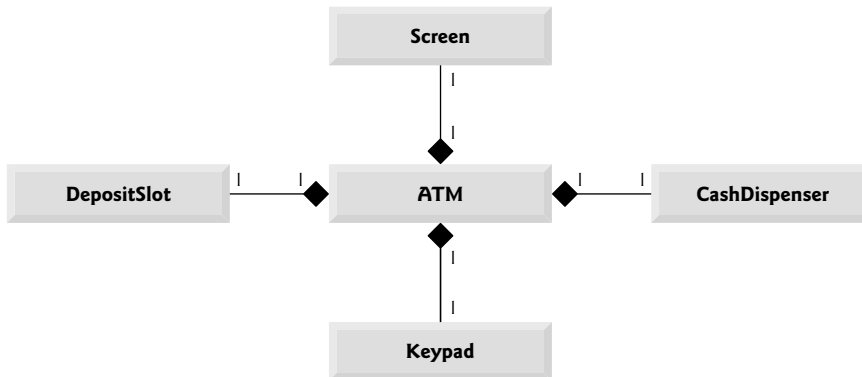


Fig. 30.9 | Class diagram showing composition relationships.

According to the UML specification, composition relationships have the following properties:

1. Only one class in the relationship can represent the whole (i.e., the diamond can be placed on only one end of the association line). For example, either the screen is part of the ATM or the ATM is part of the screen, but the screen and the ATM cannot both represent the whole in the relationship.

2. The parts in the composition relationship exist only as long as the whole, and the whole is responsible for creating and destroying its parts. For example, the act of constructing an ATM includes manufacturing its parts. Furthermore, if the ATM is destroyed, its screen, keypad, cash dispenser and deposit slot are also destroyed.
3. A part may belong to only one whole at a time, although the part may be removed and attached to another whole, which then assumes responsibility for the part.

The solid diamonds in our class diagrams indicate composition relationships that fulfill these three properties. If a *has-a* relationship does not satisfy one or more of these criteria, the UML specifies that hollow diamonds be attached to the ends of association lines to indicate **aggregation**—a weaker form of composition. For example, a personal computer and a computer monitor participate in an aggregation relationship—the computer “has a” monitor, but the two parts can exist independently, and the same monitor can be attached to multiple computers at once, thus violating the second and third properties of composition.

Figure 30.10 shows a class diagram for the ATM system. This diagram models most of the classes that we identified earlier in this section, as well as the associations between them that we can infer from the requirements document. [Note: Classes *BalanceInquiry* and *Deposit* participate in associations similar to those of class *Withdrawal*, so we’ve chosen to omit them from this diagram for simplicity. In Section 31.3, we expand our class diagram to include all the classes in the ATM system.]

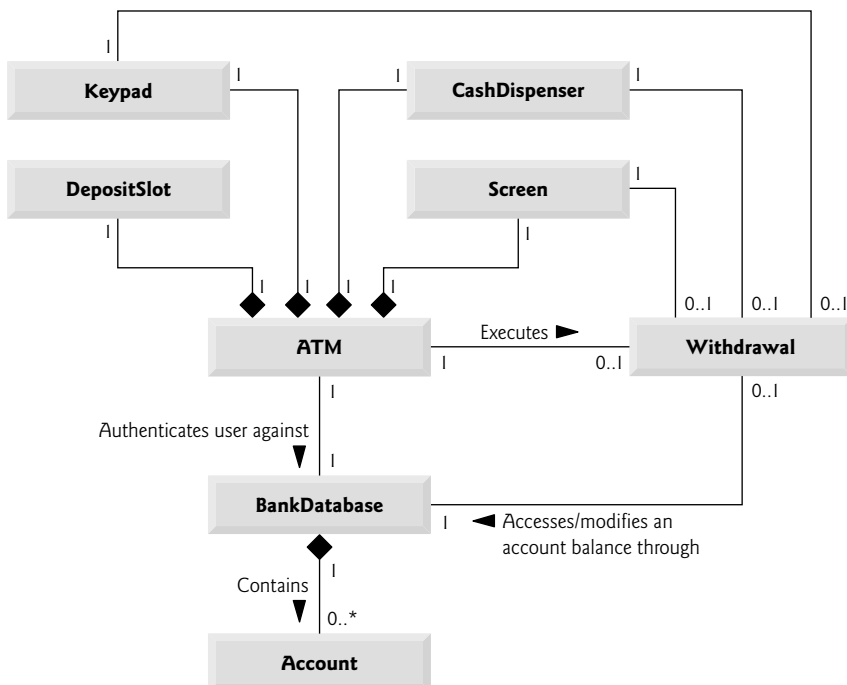


Fig. 30.10 | Class diagram for the ATM system model.

Figure 30.10 presents a graphical model of the structure of the ATM system. This class diagram includes classes `BankDatabase` and `Account` and several associations that were not present in either Fig. 30.7 or Fig. 30.9. The class diagram shows that class `ATM` has a **one-to-one relationship** with class `BankDatabase`—one `ATM` object authenticates users against one `BankDatabase` object. In Fig. 30.10, we also model the fact that the bank’s database contains information about many accounts—one object of class `BankDatabase` participates in a composition relationship with zero or more objects of class `Account`. Recall from Fig. 30.8 that the multiplicity value `0..*` at the `Account` end of the association between class `BankDatabase` and class `Account` indicates that zero or more objects of class `Account` take part in the association. Class `BankDatabase` has a **one-to-many relationship** with class `Account`—the `BankDatabase` can contain many `Accounts`. Similarly, class `Account` has a **many-to-one relationship** with class `BankDatabase`—there can be many `Accounts` in the `BankDatabase`. Recall from Fig. 30.8 that the multiplicity value `*` is identical to `0`.]

Figure 30.10 also indicates that if the user is performing a withdrawal, “one object of class `Withdrawal` accesses/modifies an account balance through one object of class `BankDatabase`.” We could have created an association directly between class `Withdrawal` and class `Account`. The requirements document, however, states that the “ATM must interact with the bank’s account-information database” to perform transactions. A bank account contains sensitive information, and systems engineers must always consider the security of personal data when designing a system. Thus, only the `BankDatabase` can access and manipulate an account directly. All other parts of the system must interact with the database to retrieve or update account information (e.g., an account balance).

The class diagram in Fig. 30.10 also models associations between class `Withdrawal` and classes `Screen`, `CashDispenser` and `Keypad`. A withdrawal transaction includes prompting the user to choose a withdrawal amount and receiving numeric input. These actions require the use of the screen and the keypad, respectively. Dispensing cash to the user requires access to the cash dispenser.

Classes `BalanceInquiry` and `Deposit`, though not shown in Fig. 30.10, take part in several associations with the other classes of the ATM system. Like class `Withdrawal`, each of these classes associates with classes `ATM` and `BankDatabase`. An object of class `BalanceInquiry` also associates with an object of class `Screen` to display the balance of an account to the user. Class `Deposit` associates with classes `Screen`, `Keypad` and `DepositSlot`. Like withdrawals, deposit transactions require use of the screen and the keypad to display prompts and receive inputs, respectively. To receive a deposit envelope, an object of class `Deposit` associates with an object of class `DepositSlot`.

We’ve identified our ATM system’s classes, although we may discover others as we proceed with the design and implementation. In Section 30.4, we determine each class’s attributes, and in Section 30.5, we use these attributes to examine how the system changes over time.

Self-Review Exercises

30.4 Suppose we have a class `Car` that represents a car. Think of some of the different pieces that a manufacturer would put together to produce a whole car. Create a class diagram (similar to Fig. 30.9) that models some of the composition relationships of class `Car`.

30.5 Suppose we have a class `File` that represents an electronic document in a stand-alone, non-networked computer represented by class `Computer`. What sort of association exists between class `Computer` and class `File`?

- a) Class `Computer` has a one-to-one relationship with class `File`.
- b) Class `Computer` has a many-to-one relationship with class `File`.
- c) Class `Computer` has a one-to-many relationship with class `File`.
- d) Class `Computer` has a many-to-many relationship with class `File`.

30.6 State whether the following statement is *true* or *false*. If *false*, explain why: A UML class diagram in which a class's second and third compartments are not modeled is said to be an elided diagram.

30.7 Modify the class diagram of Fig. 30.10 to include class `Deposit` instead of class `Withdrawal`.

30.4 Identifying Class Attributes

In the previous section, we began the first stage of an object-oriented design (OOD) for our ATM system—analyzing the requirements document and identifying the classes needed to implement the system. We listed the nouns and noun phrases in the requirements document and identified a separate class for each one that plays a significant role in the ATM system. We then modeled the classes and their relationships in a UML class diagram (Fig. 30.10). Classes have attributes (data) and operations (behaviors). Class attributes are implemented in C# programs as instance variables and properties, and class operations are implemented as methods and properties. In this section, we determine many of the attributes needed in the ATM system. In Section 30.5, we examine how these attributes represent an object's state. In Section 30.6, we determine the operations for our classes.

Identifying Attributes

Consider the attributes of some real-world objects: A person's attributes include height, weight and whether the person is left-handed, right-handed or ambidextrous. A radio's attributes include its station setting, its volume setting and its AM or FM setting. A car's attributes include its speedometer and odometer readings, the amount of gas in its tank and what gear it is in. A personal computer's attributes include its manufacturer (e.g., Dell, Gateway, Sun, Apple or IBM), type of screen (e.g., LCD or CRT), main memory size and hard-disk size.

We can identify many attributes of the classes in our system by looking for descriptive words and phrases in the requirements document. For each one we find that plays a significant role in the ATM system, we create an attribute and assign it to one or more of the classes identified in Section 30.3. We also create attributes to represent any additional data that a class may need, as such needs become clear throughout the design process.

Figure 30.11 lists the words or phrases from the requirements document that describe each class. For example, the requirements document describes the steps taken to obtain a "withdrawal amount," so we list "amount" next to class `Withdrawal`.

Figure 30.11 leads us to create one attribute of class `ATM`. Class `ATM` maintains information about the state of the ATM. The phrase "user is authenticated" describes a state of the ATM (we discuss states in detail in Section 30.5), so we include `userAuthenticated` as a **bool attribute** (i.e., an attribute that has a value of either `true` or `false`). This attribute indicates whether the ATM has successfully authenticated the current user—`userAuthenticated` must be `true` for the system to allow the user to perform transactions and

Class	Descriptive words and phrases
ATM	user is authenticated
BalanceInquiry	account number
Withdrawal	account number amount
Deposit	account number amount
BankDatabase	[no descriptive words or phrases]
Account	account number PIN balance
Screen	[no descriptive words or phrases]
Keypad	[no descriptive words or phrases]
CashDispenser	begins each day loaded with 500 \$20 bills
DepositSlot	[no descriptive words or phrases]

Fig. 30.11 | Descriptive words and phrases from the ATM requirements document.

access account information. This attribute helps ensure the security of the data in the system.

Classes `BalanceInquiry`, `Withdrawal` and `Deposit` share one attribute. Each transaction involves an “account number” that corresponds to the account of the user making the transaction. We assign integer attribute `accountNumber` to each transaction class to identify the account to which an object of the class applies.

Descriptive words and phrases in the requirements document also suggest some differences in the attributes required by each transaction class. The requirements document indicates that to withdraw cash or deposit funds, users must enter a specific “amount” of money to be withdrawn or deposited, respectively. Thus, we assign to classes `Withdrawal` and `Deposit` an attribute `amount` to store the value supplied by the user. The amounts of money related to a withdrawal and a deposit are defining characteristics of these transactions that the system requires for them to take place. Recall that `C#` represents monetary amounts with type `decimal`. Class `BalanceInquiry` does not need additional data to perform its task—it requires only an account number to indicate the account whose balance should be retrieved.

Class `Account` has several attributes. The requirements document states that each bank account has an “account number” and a “PIN,” which the system uses for identifying accounts and authenticating users. We assign to class `Account` two integer attributes: `accountNumber` and `pin`. The requirements document also specifies that an account maintains a “balance” of the amount of money in the account, and that the money the user deposits does not become available for a withdrawal until the bank verifies the amount of cash in the deposit envelope and any checks in the envelope clear. An account must still record the amount of money that a user deposits, however. Therefore, we decide that an account should represent a balance using two `decimal` attributes—`availableBalance` and

`totalBalance`. Attribute `availableBalance` tracks the amount of money that a user can withdraw from the account. Attribute `totalBalance` refers to the total amount of money that the user has “on deposit” (i.e., the amount of money available, plus the amount of cash deposits waiting to be verified or the amount of checks waiting to be cleared). For example, suppose an ATM user deposits \$50.00 in cash into an empty account. The `totalBalance` attribute would increase to \$50.00 to record the deposit, but the `availableBalance` would remain at \$0 until a bank employee counts the amount of cash in the envelope and confirms the total. [Note: We assume that the bank updates the `availableBalance` attribute of an `Account` soon after the ATM transaction occurs, in response to confirming that \$50 worth of cash was found in the deposit envelope. We assume that this update occurs through a transaction that a bank employee performs using a bank system other than the ATM. Thus, we do not discuss this transaction in our case study.]

Class `CashDispenser` has one attribute. The requirements document states that the cash dispenser “begins each day loaded with 500 \$20 bills.” The cash dispenser must keep track of the number of bills it contains to determine whether enough cash is on hand to satisfy withdrawal requests. We assign to class `CashDispenser` integer attribute `count`, which is initially set to 500.

For real problems in industry, there is no guarantee that requirements documents will be rich enough and precise enough for the object-oriented systems designer to determine all the attributes, or even all the classes. The need for additional classes, attributes and behaviors may become clear as the design process proceeds. As we progress through this case study, we too will continue to add, modify and delete information about the classes in our system.

Modeling Attributes

The class diagram in Fig. 30.12 lists some of the attributes for the classes in our system—the descriptive words and phrases in Fig. 30.11 helped us identify these attributes. For simplicity, Fig. 30.12 does not show the associations among classes—we showed these in Fig. 30.10. Systems designers commonly do this. Recall that in the UML, a class’s attributes are placed in the middle compartment of the class’s rectangle. We list each attribute’s name and type separated by a colon (:), followed in some cases by an equal sign (=) and an initial value.

Consider the `userAuthenticated` attribute of class `ATM`:

```
userAuthenticated : bool = false
```

This attribute declaration contains three pieces of information about the attribute. The **attribute name** is `userAuthenticated`. The **attribute type** is `bool`. In C#, an attribute can be represented by a simple type, such as `bool`, `int`, `double` or `decimal`, or a class type. We have chosen to model only simple-type attributes in Fig. 30.12—we discuss the reasoning behind this decision shortly.

We can also indicate an initial value for an attribute. Attribute `userAuthenticated` in class `ATM` has an initial value of `false`. This indicates that the system initially does not consider the user to be authenticated. If an attribute has no initial value specified, only its name and type (separated by a colon) are shown. For example, the `accountNumber` attribute of class `BalanceInquiry` is an `int`. Here we show no initial value, because the value of this attribute is a number that we do not yet know. This number will be determined at execution time based on the account number entered by the current ATM user.

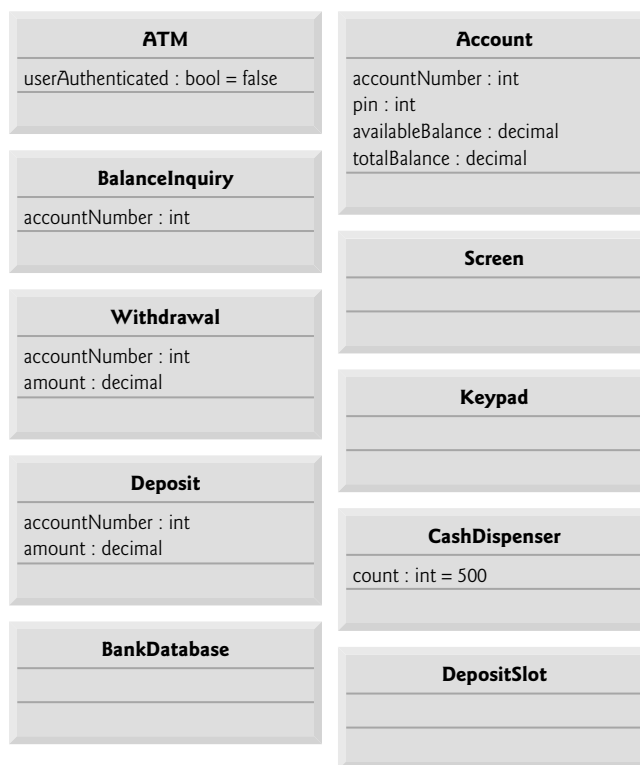


Fig. 30.12 | Classes with attributes.

Figure 30.12 does not contain attributes for classes `Screen`, `Keypad` and `DepositSlot`. These are important components of our system for which our design process simply has not yet revealed any attributes. We may discover some, however, in the remaining phases of design or when we implement these classes in C#. This is perfectly normal.



Software Engineering Observation 30.1

Early in the design process, classes often lack attributes (and operations). Such classes should not be eliminated, however, because attributes (and operations) may become evident in the later phases of design and implementation.

Fig. 30.12 also does not include attributes for class `BankDatabase`. We have chosen to include only simple-type attributes in Fig. 30.12 (and in similar class diagrams throughout the case study). A class-type attribute is modeled more clearly as an association (in particular, a composition) between the class with the attribute and the attribute's own class. For example, the class diagram in Fig. 30.10 indicates that class `BankDatabase` participates in a composition relationship with zero or more `Account` objects. From this composition, we can determine that when we implement the ATM system in C#, we'll be required to create an attribute of class `BankDatabase` to hold zero or more `Account` objects. Similarly, we'll assign attributes to class `ATM` that correspond to its composition relationships with classes `Screen`, `Keypad`, `CashDispenser` and `DepositSlot`. These composition-based attributes

would be redundant if modeled in Fig. 30.12, because the compositions modeled in Fig. 30.10 already convey the fact that the database contains information about zero or more accounts and that an ATM is composed of a screen, keypad, cash dispenser and deposit slot. Software developers typically model these whole/part relationships as composition associations rather than as attributes required to implement the relationships.

The class diagram in Fig. 30.12 provides a solid basis for the structure of our model, but the diagram is not complete. In Section 30.5 we identify the states and activities of the objects in the model, and in Section 30.6 we identify the operations that the objects perform. As we present more of the UML and object-oriented design, we'll continue to strengthen the structure of our model.

Self-Review Exercises

30.8 We typically identify the attributes of the classes in our system by analyzing the _____ in the requirements document.

- a) nouns and noun phrases
- b) descriptive words and phrases
- c) verbs and verb phrases
- d) All of the above

30.9 Which of the following is not an attribute of an airplane?

- a) length
- b) wingspan
- c) fly
- d) number of seats

30.10 Describe the meaning of the following attribute declaration of class `CashDispenser` in the class diagram in Fig. 30.12:

```
count : int = 500
```

30.5 Identifying Objects' States and Activities

In the previous section, we identified many of the class attributes needed to implement the ATM system and added them to the class diagram in Fig. 30.12. In this section, we show how these attributes represent an object's state. We identify some key states that our objects may occupy and discuss how objects change state in response to various events occurring in the system. We also discuss the workflow, or **activities**, that various objects perform in the ATM system. We present the activities of `BalanceInquiry` and `Withdrawal` transaction objects in this section.

State Machine Diagrams

Each object in a system goes through a series of discrete states. An object's state at a given point in time is indicated by the values of its attributes at that time. **State machine diagrams** model key states of an object and show under what circumstances the object changes state. Unlike the class diagrams presented in earlier case study sections, which focused primarily on the *structure* of the system, state machine diagrams model some of the *behavior* of the system.

Figure 30.13 is a simple state machine diagram that models two of the states of an object of class `ATM`. The UML represents each state in a state machine diagram as a **rounded rectangle** with the name of the state placed inside it. A **solid circle** with an

attached stick arrowhead designates the **initial state**. Recall that we modeled this state information as the `bool` attribute `userAuthenticated` in the class diagram of Fig. 30.12. This attribute is initialized to `false`, or the “User not authenticated” state, according to the state machine diagram.




Fig. 30.13 | State machine diagram for some of the states of the ATM object.

The arrows with stick arrowheads indicate **transitions** between states. An object can transition from one state to another in response to various events that occur in the system. The name or description of the event that causes a transition is written near the line that corresponds to the transition. For example, the ATM object changes from the “User not authenticated” state to the “User authenticated” state after the bank database authenticates the user. Recall from the requirements document that the database authenticates a user by comparing the account number and PIN entered by the user with those of the corresponding account in the database. If the database indicates that the user has entered a valid account number and the correct PIN, the ATM object transitions to the “User authenticated” state and changes its `userAuthenticated` attribute to the value `true`. When the user exits the system by choosing the “exit” option from the main menu, the ATM object returns to the “User not authenticated” state in preparation for the next ATM user.



Software Engineering Observation 30.2

Software designers do not generally create state machine diagrams showing every possible state and state transition for all attributes—there are simply too many of them. State machine diagrams typically show only the most important or complex states and state transitions.

Activity Diagrams

Like a state machine diagram, an activity diagram models aspects of system behavior. Unlike a state machine diagram, an activity diagram models an object’s workflow (sequence of tasks) during application execution. An activity diagram models the actions to perform and in what order the object will perform them. The activity diagram in Fig. 30.14 models the actions involved in executing a `BalanceInquiry` transaction. We assume that a `BalanceInquiry` object has already been initialized and assigned a valid account number (that of the current user), so the object knows which balance to retrieve. The diagram includes the actions that occur after the user selects a balance inquiry from the main menu and before the ATM returns the user to the main menu—a `BalanceInquiry` object does not perform or initiate these actions, so we do not model them here. The diagram begins with the retrieval of the available balance of the user’s account from the database. Next, the `BalanceInquiry` retrieves the total balance of the account. Finally, the transaction displays the balances on the screen.

The UML represents an action in an activity diagram as an action state, which is modeled by a rectangle with its left and right sides replaced by arcs curving outward. Each

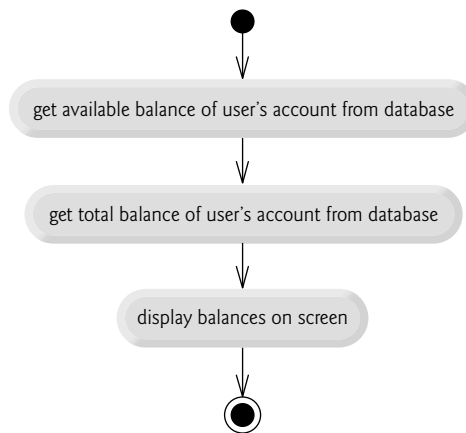


Fig. 30.14 | Activity diagram for a `BalanceInquiry` transaction.

action state contains an action expression—for example, “get available balance of user’s account from database”—that specifies an action to perform. An arrow with a stick arrow-head connects two action states, indicating the order in which the actions represented by the action states occur. The solid circle (at the top of Fig. 30.14) represents the activity’s initial state—the beginning of the workflow before the object performs the modeled actions. In this case, the transaction first executes the “get available balance of user’s account from database” action expression. Second, the transaction retrieves the total balance. Finally, the transaction displays both balances on the screen. The solid circle enclosed in an open circle (at the bottom of Fig. 30.14) represents the final state—the end of the workflow after the object performs the modeled actions.

Figure 30.15 shows an activity diagram for a `Withdrawal` transaction. We assume that a `Withdrawal` object has been assigned a valid account number. We do not model the user selecting a withdrawal from the main menu or the ATM returning the user to the main menu, because these are not actions performed by a `Withdrawal` object. The transaction first displays a menu of standard withdrawal amounts (Fig. 30.3) and an option to cancel the transaction. The transaction then inputs a menu selection from the user. The activity flow now arrives at a decision symbol. This point determines the next action based on the associated guard conditions. If the user cancels the transaction, the system displays an appropriate message. Next, the cancellation flow reaches a merge symbol, where this activity flow joins the transaction’s other possible activity flows (which we discuss shortly). A merge can have any number of incoming transition arrows, but only one outgoing transition arrow. The decision at the bottom of the diagram determines whether the transaction should repeat from the beginning. When the user has canceled the transaction, the guard condition “cash dispensed or user canceled transaction” is true, so control transitions to the activity’s final state.

If the user selects a withdrawal amount from the menu, `amount` (an attribute of class `Withdrawal` originally modeled in Fig. 30.12) is set to the value chosen by the user. The transaction next gets the available balance of the user’s account (i.e., the `availableBalance` attribute of the user’s `Account` object) from the database. The activity flow then arrives at another decision. If the requested withdrawal amount exceeds the user’s available balance, the system displays an appropriate error message informing the user of the

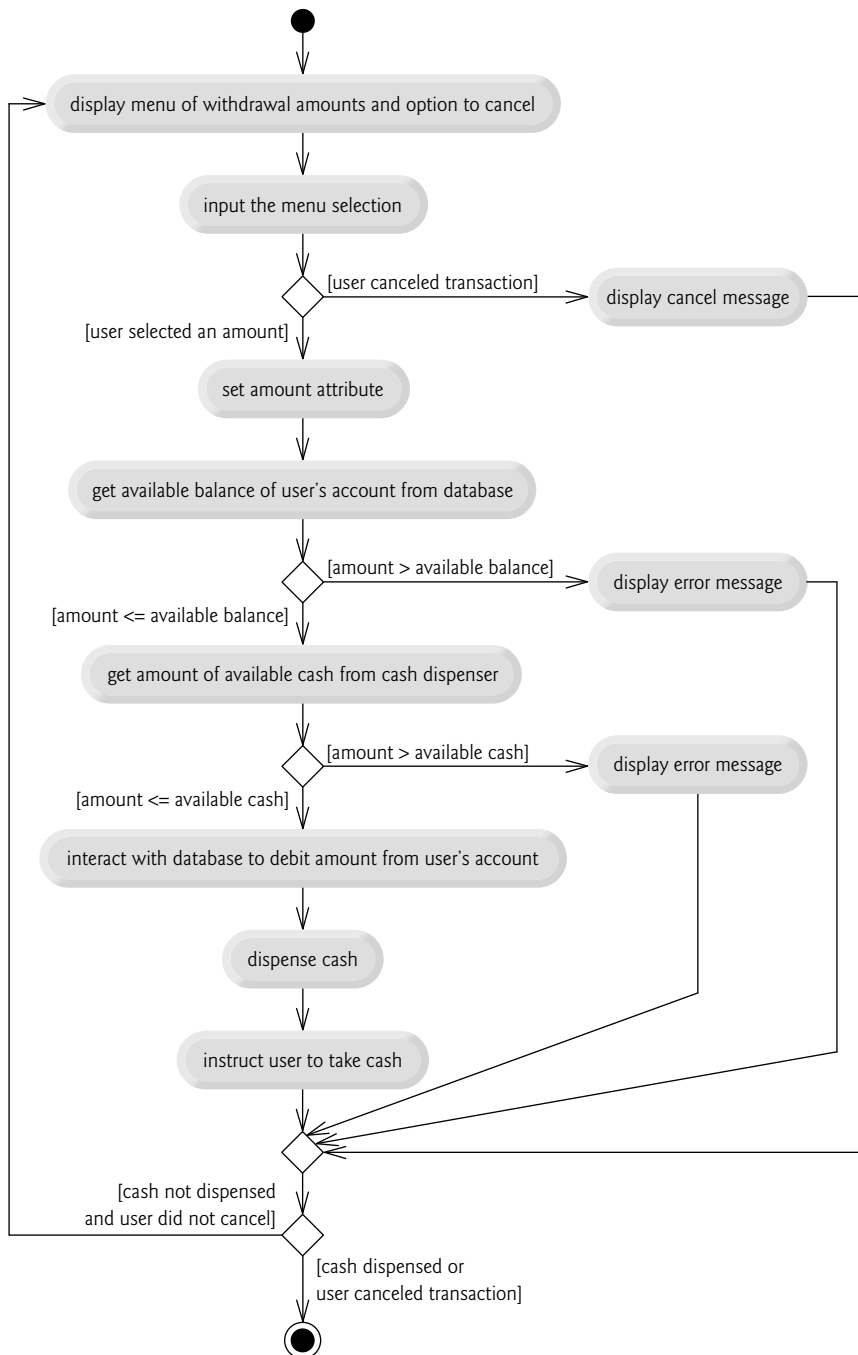


Fig. 30.15 | Activity diagram for a Withdrawal transaction.

problem. Control then merges with the other activity flows before reaching the decision at the bottom of the diagram. The guard condition “cash not dispensed and user did not cancel” is true, so the activity flow returns to the top of the diagram, and the transaction prompts the user to input a new amount.

If the requested withdrawal amount is less than or equal to the user’s available balance, the transaction tests whether the cash dispenser has enough cash to satisfy the withdrawal request. If it does not, the transaction displays an appropriate error message and passes through the merge before reaching the final decision. Cash was not dispensed, so the activity flow returns to the beginning of the activity diagram, and the transaction prompts the user to choose a new amount. If sufficient cash is available, the transaction interacts with the database to debit the withdrawal amount from the user’s account (i.e., subtract the amount from *both* the `availableBalance` and `totalBalance` attributes of the user’s `Account` object). The transaction then dispenses the desired amount of cash and instructs the user to take the cash.

The main flow of activity next merges with the two error flows and the cancellation flow. In this case, cash was dispensed, so the activity flow reaches the final state.

We’ve taken the first steps in modeling the behavior of the ATM system and have shown how an object’s attributes affect the object’s activities. In Section 30.6, we investigate the operations of our classes to create a more complete model of the system’s behavior.

Self-Review Exercises

30.11 State whether the following statement is *true* or *false*, and if *false*, explain why: State machine diagrams model structural aspects of a system.

30.12 An activity diagram models the _____ that an object performs and the order in which it performs them.

- a) actions
- b) attributes
- c) states
- d) state transitions

30.13 Based on the requirements document, create an activity diagram for a deposit transaction.

30.6 Identifying Class Operations

In the preceding sections, we performed the first few steps in the object-oriented design of our ATM system. In this section, we determine some of the class operations (or behaviors) needed to implement the ATM system.

Identifying Operations

An operation is a service that objects of a class provide to clients of the class. Consider the operations of some real-world objects. A radio’s operations include setting its station and volume (typically invoked by a person adjusting the radio’s controls). A car’s operations include accelerating (invoked by the driver pressing the accelerator pedal), decelerating (invoked by the driver pressing the brake pedal or releasing the gas pedal), turning, and shifting gears. Software objects can offer operations as well—for example, a software graphics object might offer operations for drawing a circle, drawing a line and drawing a square. A spreadsheet software object might offer operations like printing the spreadsheet,

totaling the elements in a row or column and graphing information in the spreadsheet as a bar chart or pie chart.

We can derive many of the operations of the classes in our ATM system by examining the verbs and verb phrases in the requirements document. We then relate each of these to particular classes in our system. The verbs and verb phrases in Fig. 30.16 help us determine the operations of our classes.

Class	Verbs and verb phrases
ATM	executes financial transactions
BalanceInquiry	[none in the requirements document]
Withdrawal	[none in the requirements document]
Deposit	[none in the requirements document]
BankDatabase	authenticates a user, retrieves an account balance, credits an account, debits an account
Account	retrieves an account balance, credits a deposit amount to an account, debits a withdrawal amount to an account
Screen	displays a message to the user
Keypad	receives numeric input from the user
CashDispenser	dispenses cash, indicates whether it contains enough cash to satisfy a withdrawal request
DepositSlot	receives a deposit envelope

Fig. 30.16 | Verbs and verb phrases for each class in the ATM system.

Modeling Operations

To identify operations, we examine the verb phrases listed for each class in Fig. 30.16. The “executes financial transactions” phrase associated with class ATM implies that class ATM instructs transactions to execute. Therefore, classes *BalanceInquiry*, *Withdrawal* and *Deposit* each need an operation to provide this service to the ATM. We place this operation (which we have named *Execute*) in the third compartment of the three transaction classes in the updated class diagram of Fig. 30.17. During an ATM session, the ATM object will invoke the *Execute* operation of each transaction object to tell it to execute.

The UML represents operations (which are implemented as methods in C#) by listing the operation name, followed by a comma-separated list of parameters in parentheses, a colon and the return type:

```
operationName( parameter1, parameter2, ..., parameterN ) : returnType
```

Each parameter in the comma-separated parameter list consists of a parameter name, followed by a colon and the parameter type:

```
parameterName : parameterType
```

For the moment, we do not list the parameters of our operations—we’ll identify and model the parameters of some of the operations shortly. For some of the operations, we do not yet know the return types, so we also omit them from the diagram. These omissions

are perfectly normal at this point. As our design and implementation proceed, we'll add the remaining return types.

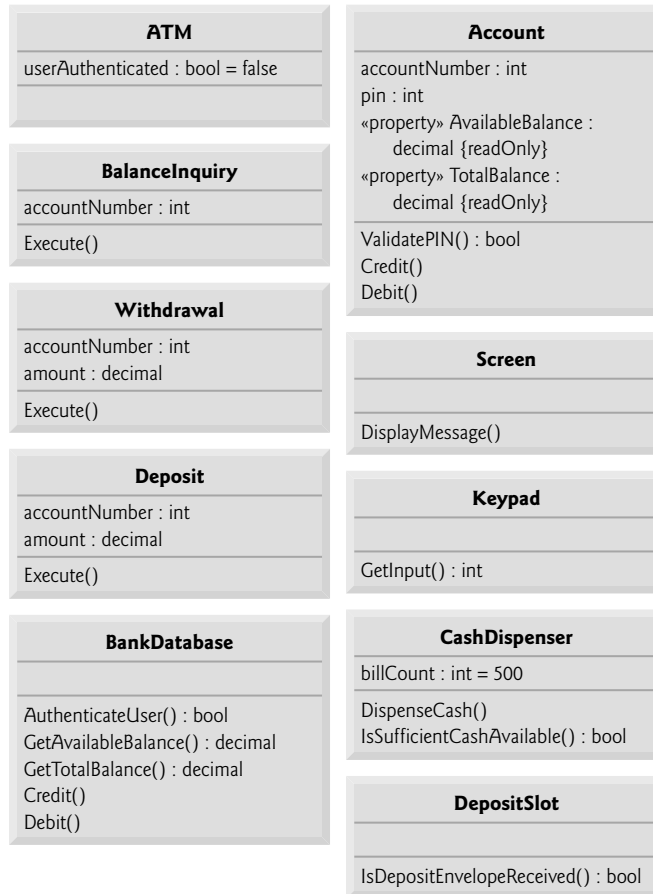


Fig. 30.17 | Classes in the ATM system with attributes and operations.

*Operations of Class **BankDatabase** and Class **Account***

Figure 30.16 lists the phrase “authenticates a user” next to class **BankDatabase**—the database is the object that contains the account information necessary to determine whether the account number and PIN entered by a user match those of an account at the bank. Therefore, class **BankDatabase** needs an operation that provides an authentication service to the ATM. We place the operation `AuthenticateUser` in the third compartment of class **BankDatabase** (Fig. 30.17). However, an object of class **Account**, not class **BankDatabase**, stores the account number and PIN that must be accessed to authenticate a user, so class **Account** must provide a service to validate a PIN obtained through user input against a PIN stored in an **Account** object. Therefore, we add a `ValidatePIN` operation to class **Account**. We specify a return type of `bool` for the `AuthenticateUser` and `ValidatePIN` operations. Each operation returns a value indicating either that the operation was successful

in performing its task (i.e., a return value of `true`) or that it was not successful (i.e., a return value of `false`).

Figure 30.16 lists several additional verb phrases for class `BankDatabase`: “retrieves an account balance,” “credits an account” and “debits an account.” Like “authenticates a user,” these remaining phrases refer to services that the database must provide to the ATM, because the database holds all the account data used to authenticate a user and perform ATM transactions. However, objects of class `Account` actually perform the operations to which these phrases refer. Thus, class `BankDatabase` and class `Account` both need operations that correspond to each of these phrases. Recall from Section 30.3 that, because a bank account contains sensitive information, we do not allow the ATM to access accounts directly. The database acts as an intermediary between the ATM and the account data, preventing unauthorized access. As we’ll see in Section 30.7, class `ATM` invokes the operations of class `BankDatabase`, each of which in turn invokes corresponding operations (which are get accessors of read-only properties) in class `Account`.

The phrase “retrieves an account balance” suggests that classes `BankDatabase` and `Account` each need an operation that gets the balance. However, recall that we created two attributes in class `Account` to represent a balance—`availableBalance` and `totalBalance`. A balance inquiry requires access to both balance attributes so that it can display them to the user, but a withdrawal needs to check only the value of `availableBalance`. To allow objects in the system to obtain these balance attributes individually from a specific `Account` object in the `BankDatabase`, we add operations `GetAvailableBalance` and `GetTotalBalance` to the third compartment of class `BankDatabase` (Fig. 30.17). We specify a return type of `decimal` for each of these operations, because the balances that they retrieve are of type `decimal`.

Once the `BankDatabase` knows which `Account` to access, it must be able to obtain each balance attribute individually from that `Account`. For this purpose, we could add operations `GetAvailableBalance` and `GetTotalBalance` to the third compartment of class `Account` (Fig. 30.17). However, in C#, simple operations such as getting the value of an attribute are typically performed by a property’s get accessor (at least when that particular class “owns” the underlying attribute). This design is for a C# application, so, rather than modeling operations `GetAvailableBalance` and `GetTotalBalance`, we model `decimal` properties `AvailableBalance` and `TotalBalance` in class `Account`. Properties are placed in the second compartment of a class diagram. These properties replace the `availableBalance` and `totalBalance` attributes that we modeled for class `Account` previously. Recall that a property’s accessors are implied—thus, they’re not modeled in a class diagram. Figure 30.16 does not mention the need to set the balances, so Fig. 30.17 shows properties `AvailableBalance` and `TotalBalance` as read-only properties (i.e., they have only get accessors). To indicate a read-only property in the UML, we follow the property’s type with “{readOnly}.”

You may be wondering why we modeled `AvailableBalance` and `TotalBalance` *properties* in class `Account`, but modeled `GetAvailableBalance` and `GetTotalBalance` *operations* in class `BankDatabase`. Since there can be many `Account` objects in the `BankDatabase`, the ATM must specify which `Account` to access when invoking `BankDatabase` operations `GetAvailableBalance` and `GetTotalBalance`. The ATM does this by passing an account-number argument to each `BankDatabase` operation. The get accessors of the properties you’ve seen in C# code cannot receive arguments. Thus, we modeled

`GetAvailableBalance` and `GetTotalBalance` as operations in class `BankDatabase` so that we could specify parameters to which the ATM can pass arguments. Also, the underlying balance attributes are not owned by the `BankDatabase`, so `get` accessors are not appropriate here. We discuss the parameters for the `BankDatabase` operations shortly.

The phrases “credits an account” and “debits from an account” indicate that classes `BankDatabase` and `Account` must perform operations to update an account during deposits and withdrawals, respectively. We therefore assign `Credit` and `Debit` operations to classes `BankDatabase` and `Account`. You may recall that crediting an account (as in a deposit) adds an amount only to the `Account`’s total balance. Debiting an account (as in a withdrawal), on the other hand, subtracts the amount from both the total and available balances. We hide these implementation details inside class `Account`. This is a good example of encapsulation and information hiding.

If this were a real ATM system, classes `BankDatabase` and `Account` would also provide a set of operations to allow another banking system to update a user’s account balance after either confirming or rejecting all or part of a deposit. Operation `ConfirmDepositAmount`, for example, would add an amount to the `Account`’s available balance, thus making deposited funds available for withdrawal. Operation `RejectDepositAmount` would subtract an amount from the `Account`’s total balance to indicate that a specified amount, which had recently been deposited through the ATM and added to the `Account`’s total balance, was invalidated (or checks may have “bounced”). The bank would invoke operation `RejectDepositAmount` after determining either that the user failed to include the correct amount of cash or that any checks did not clear (i.e., they “bounced”). While adding these operations would make our system more complete, we do not include them in our class diagrams or implementation because they’re beyond the scope of the case study.

Operations of Class Screen

Class `Screen` “displays a message to the user” at various times in an ATM session. All visual output occurs through the screen of the ATM. The requirements document describes many types of messages (e.g., a welcome message, an error message, a thank-you message) that the screen displays to the user. The requirements document also indicates that the screen displays prompts and menus to the user. However, a prompt is really just a message describing what the user should input next, and a menu is essentially a type of prompt consisting of a series of messages (i.e., menu options) displayed consecutively. Therefore, rather than provide class `Screen` with an individual operation to display each type of message, prompt and menu, we simply create one operation that can display any message specified by a parameter. We place this operation (`DisplayMessage`) in the third compartment of class `Screen` in our class diagram (Fig. 30.17). We do not worry about the parameter of this operation at this time—we model the parameter momentarily.

Operations of Class Keypad

From the phrase “receives numeric input from the user” listed by class `Keypad` in Fig. 30.16, we conclude that class `Keypad` should perform a `GetInput` operation. Because the ATM’s keypad, unlike a computer keyboard, contains only the numbers 0–9, we specify that this operation returns an integer value. Recall from the requirements document that in different situations, the user may be required to enter a different type of number (e.g., an account number, a PIN, the number of a menu option, a deposit amount as a number of cents). Class `Keypad` simply obtains a numeric value for a client of the class—

it does not determine whether the value meets any specific criteria. Any class that uses this operation must verify that the user entered appropriate numbers and, if not, display error messages via class `Screen`. [Note: When we implement the system, we simulate the ATM's keypad with a computer keyboard, and for simplicity, we assume that the user does not enter nonnumeric input using keys on the computer keyboard that do not appear on the ATM's keypad.]

Operations of Class `CashDispenser` and Class `DepositSlot`

Figure 30.16 lists “dispenses cash” for class `CashDispenser`. Therefore, we create operation `DispenseCash` and list it under class `CashDispenser` in Fig. 30.17. Class `CashDispenser` also “indicates whether it contains enough cash to satisfy a withdrawal request.” Thus, we include `IsSufficientCashAvailable`, an operation that returns a value of type `bool`, in class `CashDispenser`. Figure 30.16 also lists “receives a deposit envelope” for class `DepositSlot`. The deposit slot must indicate whether it received an envelope, so we place the operation `IsDepositEnvelopeReceived`, which returns a `bool` value, in the third compartment of class `DepositSlot`. [Note: A real hardware deposit slot would most likely send the ATM a signal to indicate that an envelope was received. We simulate this behavior, however, with an operation in class `DepositSlot` that class `ATM` can invoke to find out whether the deposit slot received an envelope.]

Operations of Class `ATM`

We do not list any operations for class `ATM` at this time. We're not yet aware of any services that class `ATM` provides to other classes in the system. When we implement the system in C#, however, operations of this class, and additional operations of the other classes in the system, may become apparent.

Identifying and Modeling Operation Parameters

So far, we have not been concerned with the parameters of our operations—we have attempted to gain only a basic understanding of the operations of each class. Let's now take a closer look at some operation parameters. We identify an operation's parameters by examining what data the operation requires to perform its assigned task.

Consider the `AuthenticateUser` operation of class `BankDatabase`. To authenticate a user, this operation must know the account number and PIN supplied by the user. Thus we specify that operation `AuthenticateUser` takes `int` parameters `userAccountNumber` and `userPIN`, which the operation must compare to the account number and PIN of an `Account` object in the database. We prefix these parameter names with `user` to avoid confusion between the operation's parameter names and the attribute names that belong to class `Account`. We list these parameters in the class diagram in Fig. 30.18, which models only class `BankDatabase`. [Note: It is perfectly normal to model only one class in a class diagram. In this case, we're most concerned with examining the parameters of this particular class, so we omit the other classes. In class diagrams later in the case study, parameters are no longer the focus of our attention, so we omit the parameters to save space. Remember, however, that the operations listed in these diagrams still have parameters.]

Recall that the UML models each parameter in an operation's comma-separated parameter list by listing the parameter name, followed by a colon and the parameter type. Figure 30.18 thus specifies, for example, that operation `AuthenticateUser` takes two parameters—`userAccountNumber` and `userPIN`, both of type `int`.




Fig. 30.18 | Class `BankDatabase` with operation parameters.

Class `BankDatabase` operations `GetAvailableBalance`, `GetTotalBalance`, `Credit` and `Debit` also each require a `userAccountNumber` parameter to identify the account to which the database must apply the operations, so we include these parameters in the class diagram. In addition, operations `Credit` and `Debit` each require a `decimal` parameter amount to specify the amount of money to be credited or debited, respectively.

The class diagram in Fig. 30.19 models the parameters of class `Account`'s operations. Operation `ValidatePIN` requires only a `userPIN` parameter, which contains the user-specified PIN to be compared with the PIN associated with the account. Like their counterparts in class `BankDatabase`, operations `Credit` and `Debit` in class `Account` each require a `decimal` parameter amount that indicates the amount of money involved in the operation. Class `Account`'s operations do not require an account-number parameter—each can be invoked only on the `Account` object in which they're executing, so including a parameter to specify an `Account` is unnecessary.

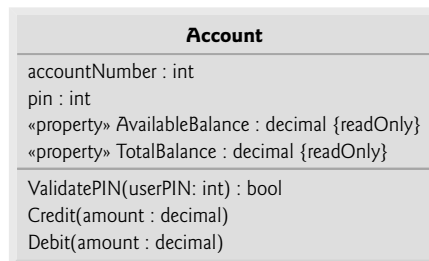


Fig. 30.19 | Class `Account` with operation parameters.

Figure 30.20 models class `Screen` with a parameter for operation `DisplayMessage`. This operation requires only `string` parameter `message`, which is the text to be displayed.




Fig. 30.20 | Class `Screen` with an operation parameter.

The class diagram in Fig. 30.21 specifies that operation `DispenseCash` of class `CashDispenser` takes decimal parameter `amount` to indicate the amount of cash (in dollars) to be dispensed. Operation `IsSufficientCashAvailable` also takes decimal parameter `amount` to indicate the amount of cash in question.




Fig. 30.21 | Class `CashDispenser` with operation parameters.

We don't discuss parameters for operation `Execute` of classes `BalanceInquiry`, `Withdrawal` and `Deposit`, operation `GetInput` of class `Keypad` and operation `IsDepositEnvelopeReceived` of class `DepositSlot`. At this point in our design process, we cannot determine whether these operations require additional data to perform their tasks, so we leave their parameter lists empty. As we progress through the case study, we may decide to add parameters to these operations.

In this section, we have determined many of the operations performed by the classes in the ATM system. We have identified the parameters and return types of some of the operations. As we continue our design process, the number of operations belonging to each class may vary—we might find that new operations are needed or that some current operations are unnecessary—and we might determine that some of our class operations need additional parameters and different return types. Again, all of this is perfectly normal.

Self-Review Exercises

30.14 Which of the following is not a behavior?

- a) reading data from a file
- b) displaying output
- c) text output
- d) obtaining input from the user

30.15 If you were to add to the ATM system an operation that returns the `amount` attribute of class `Withdrawal`, how and where would you specify this operation in the class diagram of Fig. 30.17?

30.16 Describe the meaning of the following operation listing that might appear in a class diagram for an object-oriented design of a calculator:

`Add(x : int, y : int) : int`

30.7 Identifying Collaboration Among Objects

When two objects communicate with each other to accomplish a task, they're said to **collaborate**. A **collaboration** consists of an object of one class sending a **message** to an object of another class. Messages are sent in C# via method calls. In this section, we concentrate on the collaborations (interactions) among the objects in our ATM system.

In the previous section, we determined many of the operations of the classes in our system. In this section, we concentrate on the messages that invoke these operations. To identify the collaborations in the system, we return to the requirements document of Section 30.2. Recall that this document specifies the activities that occur during an ATM session (e.g., authenticating a user, performing transactions). The steps used to describe how the system must perform each of these tasks are our first indication of the collaborations in our system. As we proceed through this and the remaining Software Engineering Case Study sections, we may discover additional collaborations.

Identifying the Collaborations in a System

We begin to identify the collaborations in the system by carefully reading the sections of the requirements document that specify what the ATM should do to authenticate a user and to perform each transaction type. For each action or step described in the requirements document, we decide which objects in our system must interact to achieve the desired result. We identify one object as the sending object (i.e., the object that sends the message) and another as the receiving object (i.e., the object that offers that operation to clients of the class). We then select one of the receiving object's operations (identified in Section 30.6) that must be invoked by the sending object to produce the proper behavior. For example, the ATM displays a welcome message when idle. We know that an object of class `Screen` displays a message to the user via its `DisplayMessage` operation. Thus, we decide that the system can display a welcome message by employing a collaboration between the ATM and the Screen in which the ATM sends a `DisplayMessage` message to the Screen by invoking the `DisplayMessage` operation of class `Screen`. [Note: To avoid repeating the phrase “an object of class...,” we refer to each object simply by using its class name preceded by an article (e.g., “a,” “an” or “the”)—for example, “the ATM” refers to an object of class `ATM`.]

Figure 30.22 lists the collaborations that can be derived from the requirements document. For each sending object, we list the collaborations in the order in which they're discussed in the requirements document. We list each collaboration involving a unique sender, message and recipient only once, even though the collaboration may occur several times during an ATM session. For example, the first row in Fig. 30.22 indicates that the ATM collaborates with the Screen whenever the ATM needs to display a message to the user.

An object of class...	sends the message...	to an object of class...
ATM	DisplayMessage GetInput AuthenticateUser Execute Execute Execute	Screen Keypad BankDatabase BalanceInquiry Withdrawal Deposit
BalanceInquiry	GetAvailableBalance GetTotalBalance DisplayMessage	BankDatabase BankDatabase Screen

Fig. 30.22 | Collaborations in the ATM system. (Part I of 2.)

An object of class...	sends the message...	to an object of class...
Withdrawal	DisplayMessage	Screen
	GetInput	Keypad
	GetAvailableBalance	BankDatabase
	IsSufficientCashAvailable	CashDispenser
	Debit	BankDatabase
Deposit	DispenseCash	CashDispenser
	DisplayMessage	Screen
	GetInput	Keypad
	IsDepositEnvelopeReceived	DepositSlot
	Credit	BankDatabase
BankDatabase	ValidatePIN	Account
	AvailableBalance (get)	Account
	TotalBalance (get)	Account
	Debit	Account
	Credit	Account

Fig. 30.22 | Collaborations in the ATM system. (Part 2 of 2.)

Let's consider the collaborations in Fig. 30.22. Before allowing a user to perform any transactions, the ATM must prompt the user to enter an account number, then a PIN. It accomplishes each of these tasks by sending a `DisplayMessage` message to the `Screen`. Both of these actions refer to the same collaboration between the ATM and the `Screen`, which is already listed in Fig. 30.22. The ATM obtains input in response to a prompt by sending a `GetInput` message to the `Keypad`. Next the ATM must determine whether the user-specified account number and PIN match those of an account in the database. It does so by sending an `AuthenticateUser` message to the `BankDatabase`. Recall that the `BankDatabase` cannot authenticate a user directly—only the user's `Account` (i.e., the `Account` that contains the account number specified by the user) can access the user's PIN to authenticate the user. Figure 30.22 therefore lists a collaboration in which the `BankDatabase` sends a `ValidatePIN` message to an `Account`.

After the user is authenticated, the ATM displays the main menu by sending a series of `DisplayMessage` messages to the `Screen` and obtains input containing a menu selection by sending a `GetInput` message to the `Keypad`. We have already accounted for these collaborations. After the user chooses a type of transaction to perform, the ATM executes the transaction by sending an `Execute` message to an object of the appropriate transaction class (i.e., a `BalanceInquiry`, a `Withdrawal` or a `Deposit`). For example, if the user chooses to perform a balance inquiry, the ATM sends an `Execute` message to a `BalanceInquiry`.

Further examination of the requirements document reveals the collaborations involved in executing each transaction type. A `BalanceInquiry` retrieves the amount of money available in the user's account by sending a `GetAvailableBalance` message to the `BankDatabase`, which sends a `get` message to an `Account`'s `AvailableBalance` property to access the available balance. Similarly, the `BalanceInquiry` retrieves the amount of money on deposit by sending a `GetTotalBalance` message to the `BankDatabase`, which sends a `get` message to an `Account`'s `TotalBalance` property to access the total balance on deposit. To display both measures of the user's balance at the same time, the `BalanceInquiry` sends `DisplayMessage` messages to the `Screen`.

A `Withdrawal` sends `DisplayMessage` messages to the `Screen` to display a menu of standard withdrawal amounts (i.e., \$20, \$40, \$60, \$100, \$200). The `Withdrawal` sends a `GetInput` message to the `Keypad` to obtain the user's menu selection. Next, the `Withdrawal` determines whether the requested withdrawal amount is less than or equal to the user's account balance. The `Withdrawal` obtains the amount of money available in the user's account by sending a `GetAvailableBalance` message to the `BankDatabase`. The `Withdrawal` then tests whether the cash dispenser contains enough cash by sending an `IsSufficientCashAvailable` message to the `CashDispenser`. A `Withdrawal` sends a `Debit` message to the `BankDatabase` to decrease the user's account balance. The `BankDatabase` in turn sends the same message to the appropriate `Account`. Recall that debiting an `Account` decreases both the total balance and the available balance. To dispense the requested amount of cash, the `Withdrawal` sends a `DispenseCash` message to the `CashDispenser`. Finally, the `Withdrawal` sends a `DisplayMessage` message to the `Screen`, instructing the user to take the cash.

A `Deposit` responds to an `Execute` message first by sending a `DisplayMessage` message to the `Screen` to prompt the user for a deposit amount. The `Deposit` sends a `GetInput` message to the `Keypad` to obtain the user's input. The `Deposit` then sends a `DisplayMessage` message to the `Screen` to tell the user to insert a deposit envelope. To determine whether the deposit slot received an incoming deposit envelope, the `Deposit` sends an `IsDepositEnvelopeReceived` message to the `DepositSlot`. The `Deposit` updates the user's account by sending a `Credit` message to the `BankDatabase`, which subsequently sends a `Credit` message to the user's `Account`. Recall that crediting an `Account` increases the total balance but not the available balance.

Interaction Diagrams

Now that we have identified a set of possible collaborations between the objects in our ATM system, let us graphically model these interactions. The UML provides several types of **interaction diagrams** that model the behavior of a system by modeling how objects interact with one another. The **communication diagram** emphasizes *which objects* participate in collaborations. [Note: Communication diagrams were called **collaboration diagrams** in earlier versions of the UML.] Like the communication diagram, the **sequence diagram** shows collaborations among objects, but it emphasizes *when* messages are sent between objects.

Communication Diagrams

Figure 30.23 shows a communication diagram that models the ATM executing a `BalanceInquiry`. Objects are modeled in the UML as rectangles containing names in the form `objectName : ClassName`. In this example, which involves only one object of each type, we disregard the object name and list only a colon followed by the class name. Specifying the name of each object in a communication diagram is recommended when modeling multiple objects of the same type. Communicating objects are connected with solid lines, and messages are passed between objects along these lines in the direction shown by arrows with filled arrowheads. The name of the message, which appears next to the arrow, is the name of an operation (i.e., a method) belonging to the receiving object—think of the name as a service that the receiving object provides to sending objects (its “clients”).




Fig. 30.23 | Communication diagram of the ATM executing a balance inquiry.

The filled arrow in Fig. 30.23 represents a message—or **synchronous call**—in the UML and a method call in C#. This arrow indicates that the flow of control is from the sending object (the ATM) to the receiving object (a `BalanceInquiry`). Since this is a synchronous call, the sending object cannot send another message, or do anything at all, until the receiving object processes the message and returns control (and possibly a return value) to the sending object. The sender just waits. For example, in Fig. 30.23, the ATM calls method `Execute` of a `BalanceInquiry` and cannot send another message until `Execute` finishes and returns control to the ATM. [Note: If this were an **asynchronous call**, represented by a stick arrowhead, the sending object would not have to wait for the receiving object to return control—it would continue sending additional messages immediately following the asynchronous call. Such calls are beyond the scope of this book.]

Sequence of Messages in a Communication Diagram

Figure 30.24 shows a communication diagram that models the interactions among objects in the system when an object of class `BalanceInquiry` executes. We assume that the object's `accountNumber` attribute contains the account number of the current user. The collaborations in Fig. 30.24 begin after the ATM sends an `Execute` message to a `BalanceInquiry` (i.e., the interaction modeled in Fig. 30.23). The number to the left of a message name indicates the order in which the message is passed. The **sequence of messages** in a communication diagram progresses in numerical order from least to greatest. In this diagram, the numbering starts with message 1 and ends with message 3. The `BalanceInquiry` first sends a `GetAvailableBalance` message to the `BankDatabase` (message 1), then sends a `GetTotalBalance` message to the `BankDatabase` (message 2). Within the parentheses following a message name, we can specify a comma-separated list of the names of the arguments sent with the message (i.e., arguments in a C# method call)—the `BalanceInquiry` passes attribute `accountNumber` with its messages to the `BankDatabase` to indicate which `Account`'s balance information to retrieve. Recall from Fig. 30.18 that operations `GetAvailableBalance` and `GetTotalBalance` of class `BankDatabase` each require a parameter to identify an account. The `BalanceInquiry` next displays the available balance and the total balance to the user by passing a `DisplayMessage` message to the `Screen` (message 3) that includes a parameter indicating the message to be displayed.

Figure 30.24 models two additional messages passing from the `BankDatabase` to an `Account` (message 1.1 and message 2.1). To provide the ATM with the two balances of the user's `Account` (as requested by messages 1 and 2), the `BankDatabase` must send `get` messages to the `Account`'s `AvailableBalance` and `TotalBalance` properties. A message passed within the handling of another message is called a **nested message**. The UML recommends using a decimal numbering scheme to indicate nested messages. For example, message 1.1 is the first message nested in message 1—the `BankDatabase` sends the `get` message to the `Account`'s `AvailableBalance` property during `BankDatabase`'s processing of a `GetAvailableBalance` message. [Note: If the `BankDatabase` needed to pass a second nested message while processing message 1, it would be numbered 1.2.] A message may

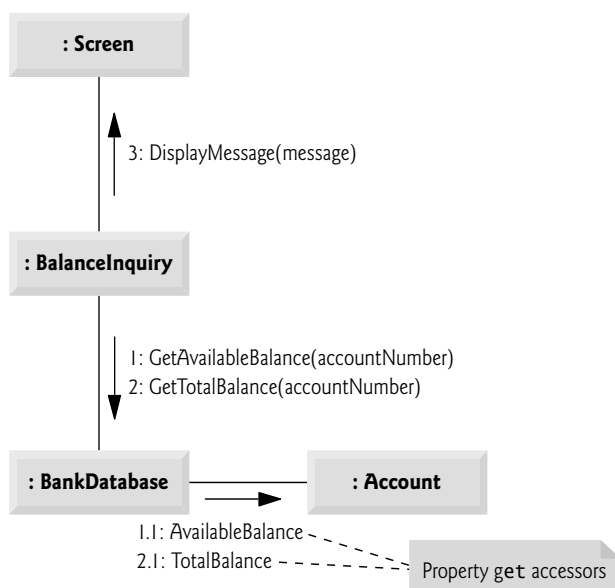


Fig. 30.24 | Communication diagram for executing a BalanceInquiry.

be passed only when all the nested messages from the previous message have been passed. For example, the BalanceInquiry passes message 3 to the Screen only after messages 2 and 2.1 have been passed, in that order.

The nested numbering scheme used in communication diagrams helps clarify precisely when and in what context each message is passed. For example, if we numbered the five messages in Fig. 30.24 using a flat numbering scheme (i.e., 1, 2, 3, 4, 5), someone looking at the diagram might not be able to determine that BankDatabase passes the get message to an Account's AvailableBalance property (message 1.1) *during* the BankDatabase's processing of message 1, as opposed to *after* completing the processing of message 1. The nested decimal numbers make it clear that the get message (message 1.1) is passed to an Account's AvailableBalance property within the handling of the GetAvailableBalance message (message 1) by the BankDatabase.

Sequence Diagrams

Communication diagrams emphasize the participants in collaborations but model their timing a bit awkwardly. A sequence diagram helps model the timing of collaborations more clearly. Figure 30.25 shows a sequence diagram modeling the sequence of interactions that occur when a Withdrawal executes. The dotted line extending down from an object's rectangle is that object's **lifeline**, which represents the progression of time. Actions typically occur along an object's lifeline in chronological order from top to bottom—an action near the top happens before one near the bottom.

Message passing in sequence diagrams is similar to message passing in communication diagrams. An arrow with a filled arrowhead extending from the sending object to the receiving object represents a message between two objects. The arrowhead points to an activation on the receiving object's lifeline. An **activation**, shown as a thin vertical rect-

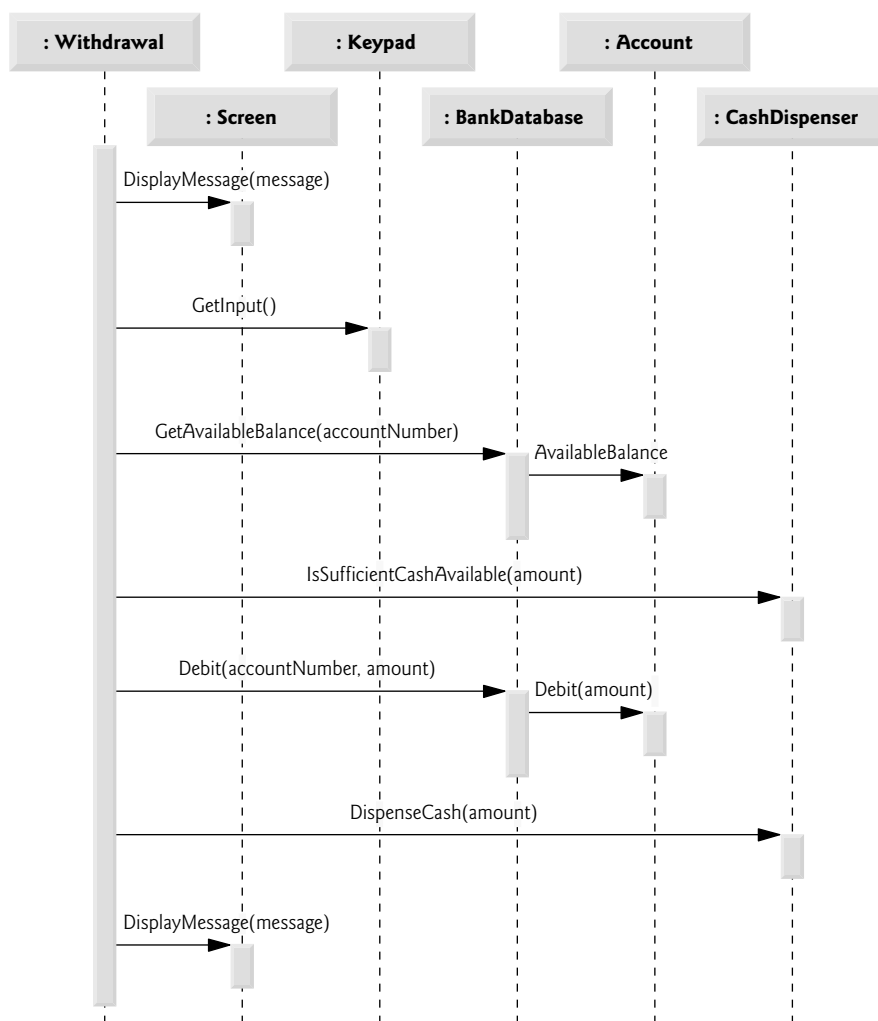


Fig. 30.25 | Sequence diagram that models a `Withdrawal` executing.

angle, indicates that an object is executing. When an object returns control, a return message, represented as a dashed line with a stick arrowhead, extends from the activation of the object returning control to the activation of the object that initially sent the message. To eliminate clutter, we omit the return-message arrows—the UML allows this practice to make diagrams more readable. Like communication diagrams, sequence diagrams can indicate message parameters between the parentheses following a message name.

The sequence of messages in Fig. 30.25 begins when a `Withdrawal` prompts the user to choose a withdrawal amount by sending a `DisplayMessage` message to the `Screen`. The `Withdrawal` then sends a `GetInput` message to the `Keypad`, which obtains input from the user. We have already modeled the control logic involved in a `Withdrawal` in the activity diagram of Fig. 30.15, so we do not show this logic in the sequence diagram of Fig. 30.25.

Instead, we model the best-case scenario, in which the balance of the user's account is greater than or equal to the chosen withdrawal amount, and the cash dispenser contains a sufficient amount of cash to satisfy the request. You can model control logic in a sequence diagram with UML frames (which are not covered in this case study). For a quick overview of UML frames, visit www.agilemodeling.com/style/frame.htm.

After obtaining a withdrawal amount, the `Withdrawal` sends a `GetAvailableBalance` message to the `BankDatabase`, which in turn sends a `get` message to the `Account's AvailableBalance` property. Assuming that the user's account has enough money available to permit the transaction, the `Withdrawal` next sends an `IsSufficientCashAvailable` message to the `CashDispenser`. Assuming that there is enough cash available, the `Withdrawal` decreases the balance of the user's account (both the total balance and the available balance) by sending a `Debit` message to the `BankDatabase`. The `BankDatabase` responds by sending a `Debit` message to the user's `Account`. Finally, the `Withdrawal` sends a `DispenseCash` message to the `CashDispenser` and a `DisplayMessage` message to the `Screen`, telling the user to remove the cash from the machine.

We have identified collaborations among objects in the ATM system and modeled some of these collaborations using UML interaction diagrams—communication diagrams and sequence diagrams. In the next section, we enhance the structure of our model to complete a preliminary object-oriented design; then we begin implementing the ATM system in C#.

Self-Review Exercises

- 30.1** A(n) _____ consists of an object of one class sending a message to an object of another class.
- association
 - aggregation
 - collaboration
 - composition
- 30.2** Which form of interaction diagram emphasizes *what* collaborations occur? Which form emphasizes *when* collaborations occur?
- 30.3** Create a sequence diagram that models the interactions among objects in the ATM system that occur when a `Deposit` executes successfully. Explain the sequence of messages modeled by the diagram.

30.8 Wrap-Up

In this chapter, you learned how to work from a detailed requirements document to develop an object-oriented design. You worked with six popular types of UML diagrams to graphically model an object-oriented automated teller machine software system. In Chapter 31, we tune the design using inheritance, then completely implement the design in a C# console application.

Answers to Self-Review Exercises

- 30.1** Figure 30.26 contains a use case diagram for a modified version of our ATM system that also allows users to transfer money between accounts.
- 30.2** b.




Fig. 30.26 | Use case diagram for a modified version of our ATM system that also allows users to transfer money between accounts.

30.3 d.

30.4 Figure 30.27 presents a class diagram that shows some of the composition relationships of a class Car.




Fig. 30.27 | Class diagram showing some composition relationships of a class Car.

30.5 c. In a computer network, this relationship could be many-to-many.

30.6 True.

30.7 Figure 30.28 presents a class diagram for the ATM including class Deposit instead of class Withdrawal (as in Fig. 30.10). Class Deposit does not associate with class CashDispenser but does associate with class DepositSlot.

30.8 b.

30.9 c. Fly is an operation or behavior of an airplane, not an attribute.

30.10 This declaration indicates that attribute count is an int with an initial value of 500; count keeps track of the number of bills available in the CashDispenser at any given time.

30.11 False. State machine diagrams model some of the behaviors of a system.

30.12 a.

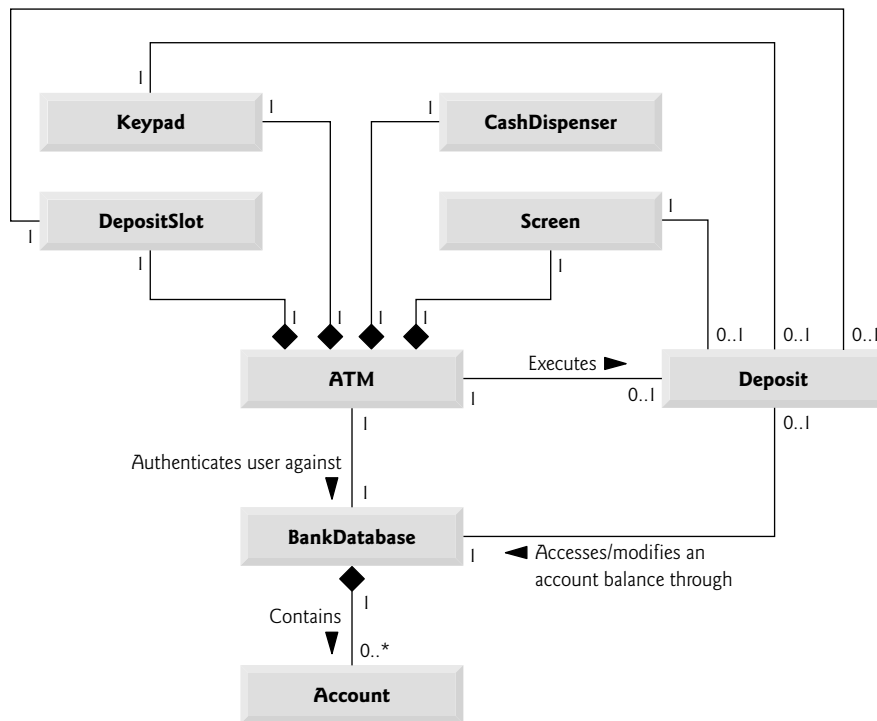


Fig. 30.28 | Class diagram for the ATM system model including class `Deposit`.

30.13 Figure 30.29 presents an activity diagram for a deposit transaction. The diagram models the actions that occur after the user chooses the deposit option from the main menu and before the ATM returns the user to the main menu. Recall that part of receiving a deposit amount from the user involves converting an integer number of cents to a dollar amount. Also recall that crediting a deposit amount to an account involves increasing only the `totalBalance` attribute of the user's `Account` object. The bank updates the `availableBalance` attribute of the user's `Account` object only after confirming the amount of cash in the deposit envelope and after the enclosed checks clear—this occurs independently of the ATM system.

30.14 c.

30.15 An operation that retrieves the amount attribute of class `Withdrawal` would typically be implemented as a get accessor of a property of class `Withdrawal`. The following would replace attribute amount in the attribute (i.e., second) compartment of class `Withdrawal`:

```
«property» Amount : decimal {readOnly}
```

30.16 This is an operation named `Add` that takes `int` parameters `x` and `y` and returns an `int` value. This operation would most likely sum its parameters `x` and `y` and return the result.

30.17 c.

30.18 Communication diagrams emphasize *what* collaborations occur. Sequence diagrams emphasize *when* collaborations occur.

30.19 Figure 30.30 presents a sequence diagram that models the interactions between objects in the ATM system that occur when a `Deposit` executes successfully. It indicates that a `Deposit` first

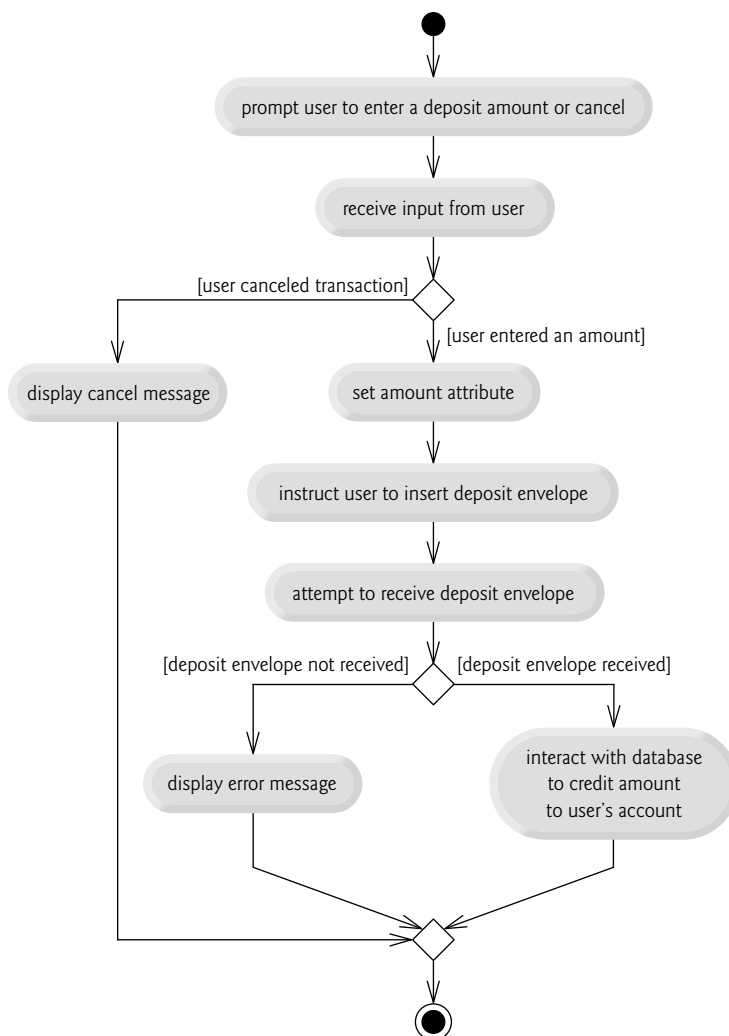


Fig. 30.29 | Activity diagram for a Deposit transaction.

sends a `DisplayMessage` message to the `Screen` (to ask the user to enter a deposit amount). Next, the `Deposit` sends a `GetInput` message to the `Keypad` to receive the amount the user will be depositing. The `Deposit` then prompts the user (to insert a deposit envelope) by sending a `DisplayMessage` message to the `Screen`. The `Deposit` next sends an `IsDepositEnvelopeReceived` message to the `DepositSlot` to confirm that the deposit envelope has been received by the ATM. Finally, the `Deposit` increases the total balance (but not the available balance) of the user's `Account` by sending a `Credit` message to the `BankDatabase`. The `BankDatabase` responds by sending the same message to the user's `Account`.

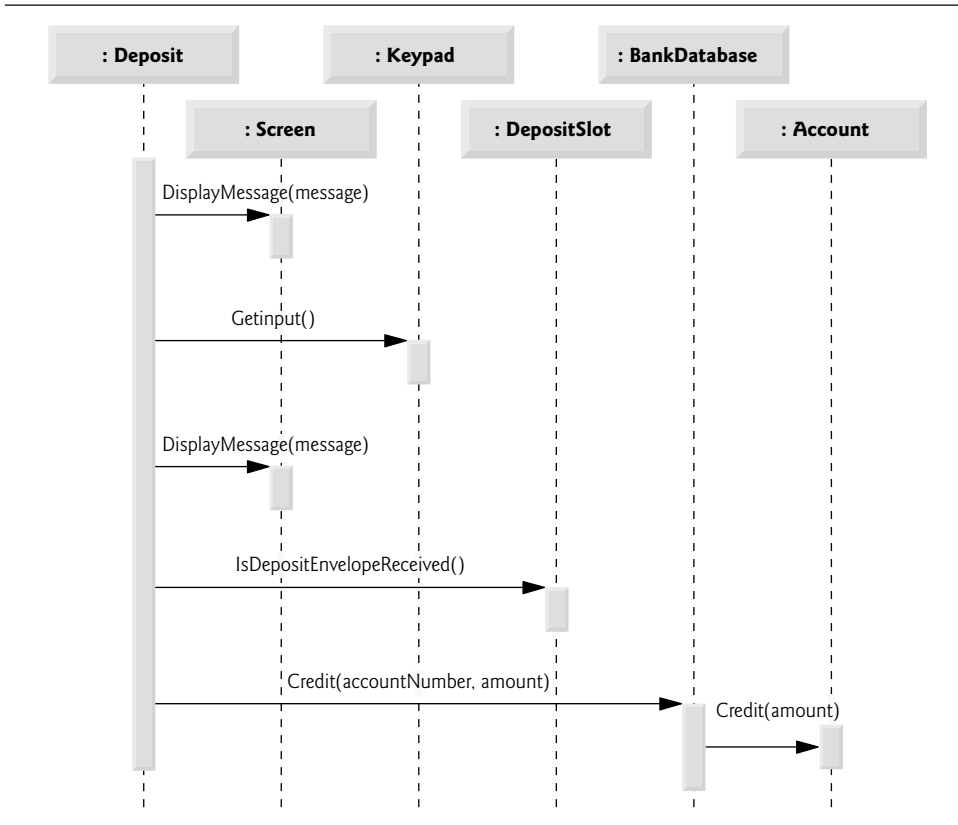


Fig. 30.30 | Sequence diagram that models a Deposit executing.

